

Network Fundamentals

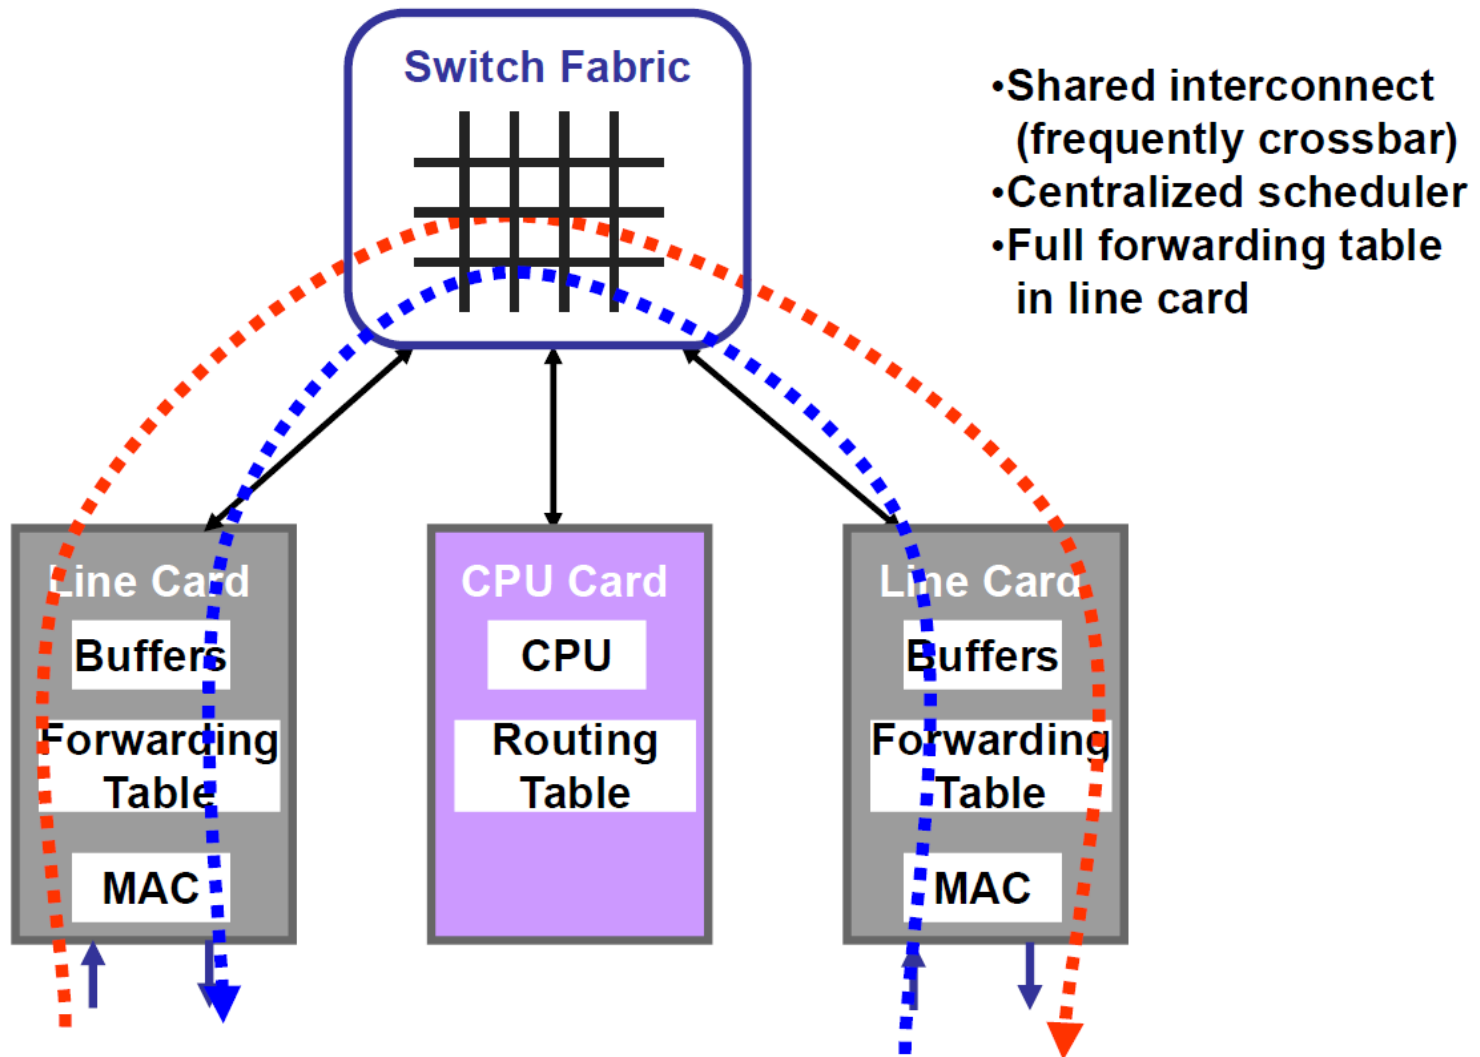
Peng Zhang

School of Computer Science and Technology


Xi'an Jiaotong University

Spring 2024


Recap: Third Generation Routers



Recap: Head of Line Blocking




Recap: Virtual output queues



Question 1

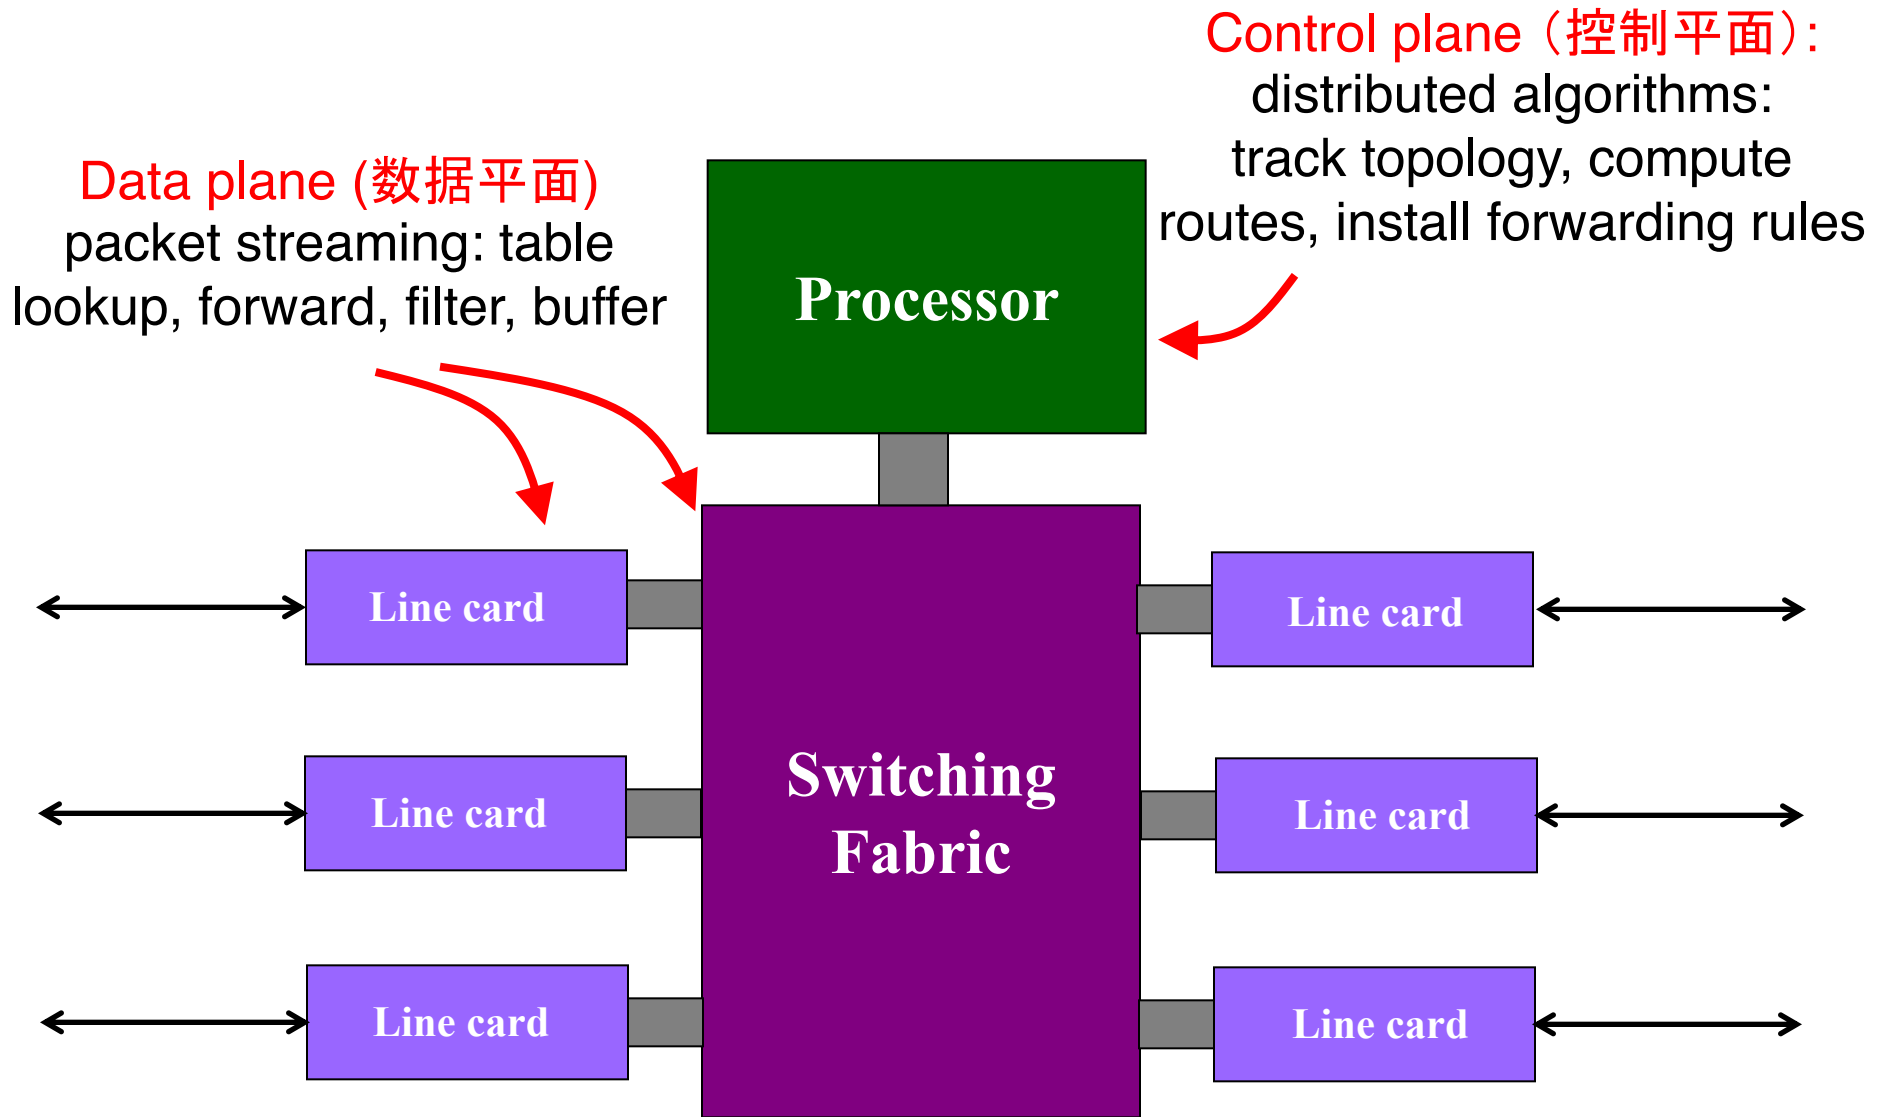
Is CSMA/CD necessary?



With the growing popularity of Ethernet switches in the 1990s, IEEE 802.3 deprecated Ethernet repeaters in 2011, making **CSMA/CD and half-duplex operation less common and less important.**



Recap: Control Plane vs. Data Plane




Today

How do **local** computers communicate?



Switching



Switching

- **MAC Address**
 - 48bits identifying the other end of link
 - Globally Unique

34:36:3b:d2:8a:86

Switching

- MAC Address is Hierarchical

34:36:3b:d2:8a:86



Apple, Inc. 1 Infinite Loop
Cupertino CA 95014 US

see <http://standards-oui.ieee.org/oui/oui.txt>

Broadcast Address

- The address with all bits set to 1 identifies the broadcast address

`ff:ff:ff:ff:ff:ff`

enables to send a frame to
all adapters on the link

What is My IP Address?

I want to send an IP packet to
192.168.1.10?

What destination MAC do I use?!



34:36:3b:d2:8a:10


Assigned by manufacture

192.168.1.9

How do I know my IP address?

DHCP

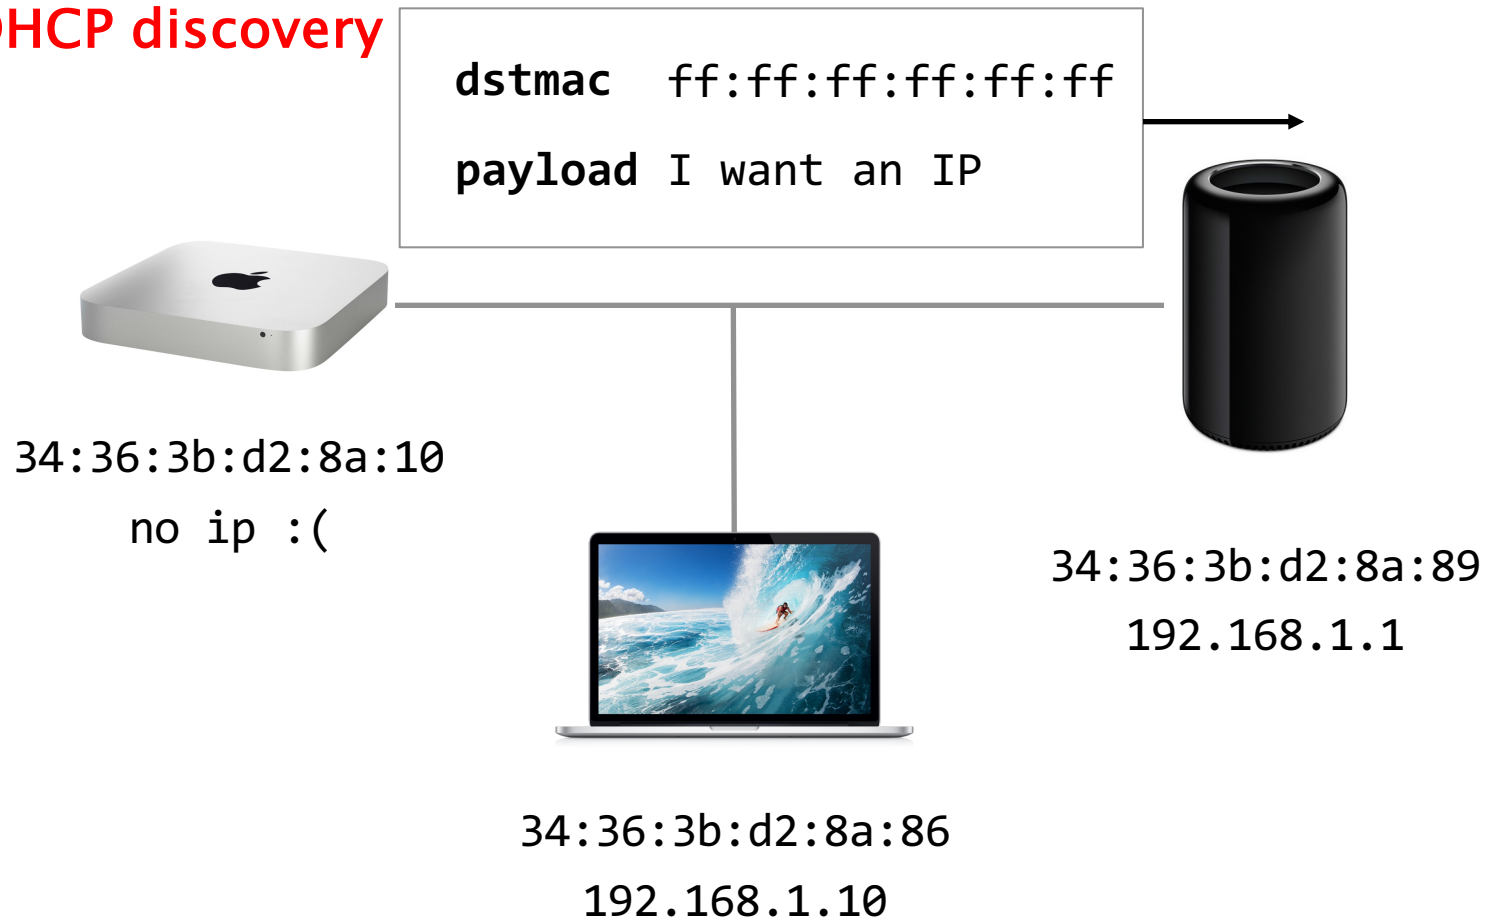
Network adapters acquire an IP address using the Dynamic Host Configuration Protocol (DHCP)



DHCP

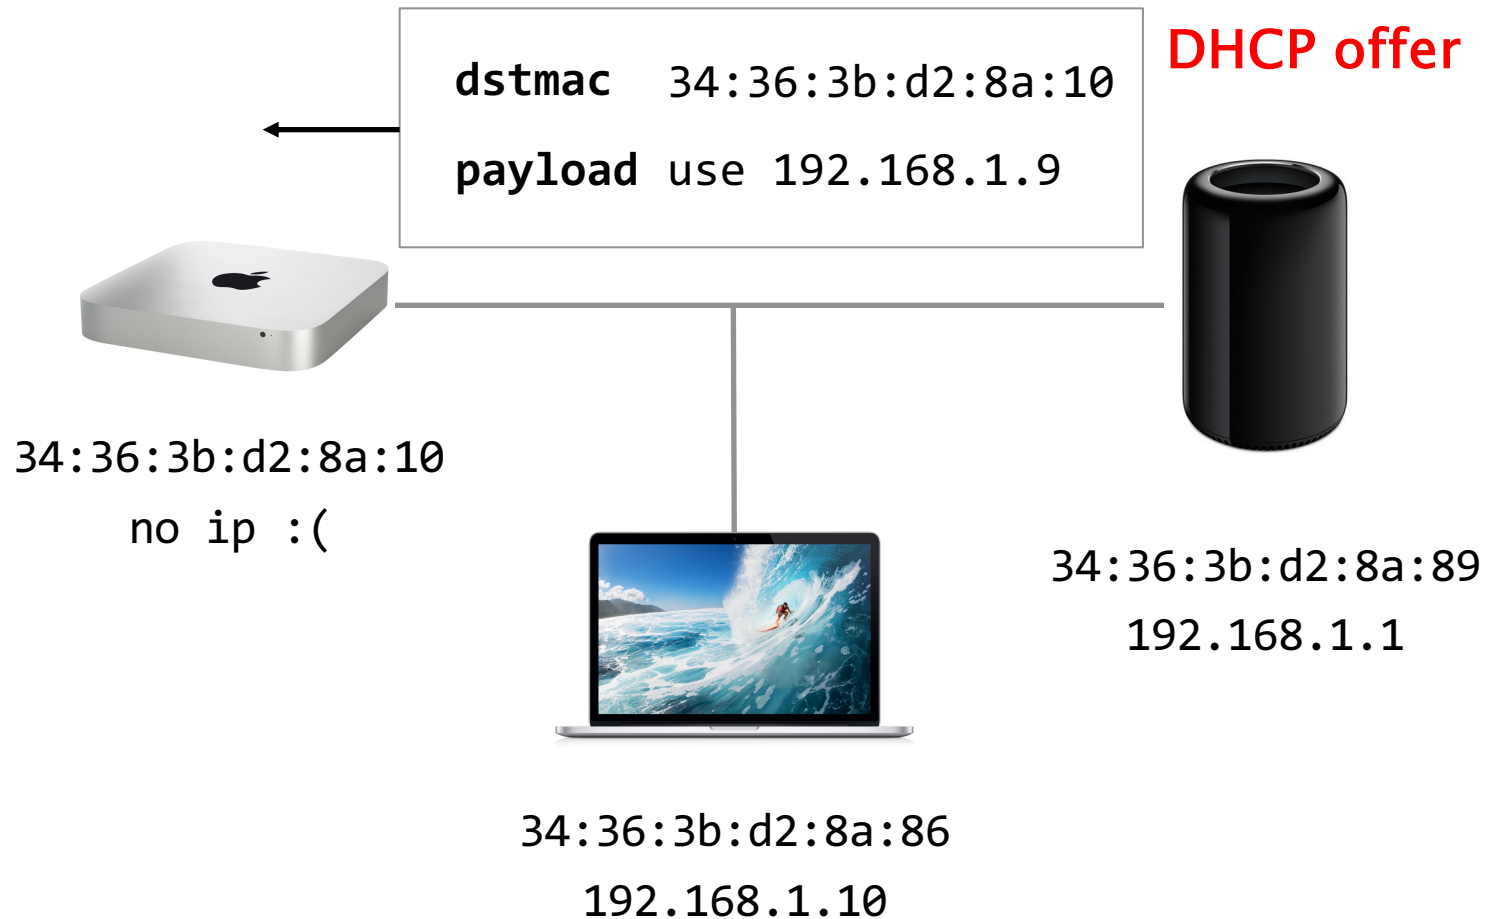
Host sends an “IP request” to everyone on the link using the broadcast address

DHCP discovery




DHCP

DHCP server (if any) answers with an IP address



DHCP



Address Resolution Protocol (ARP)

I want to send an IP packet to
192.168.1.10?

What destination MAC do I use?!




34:36:3b:d2:8a:10

192.168.1.9

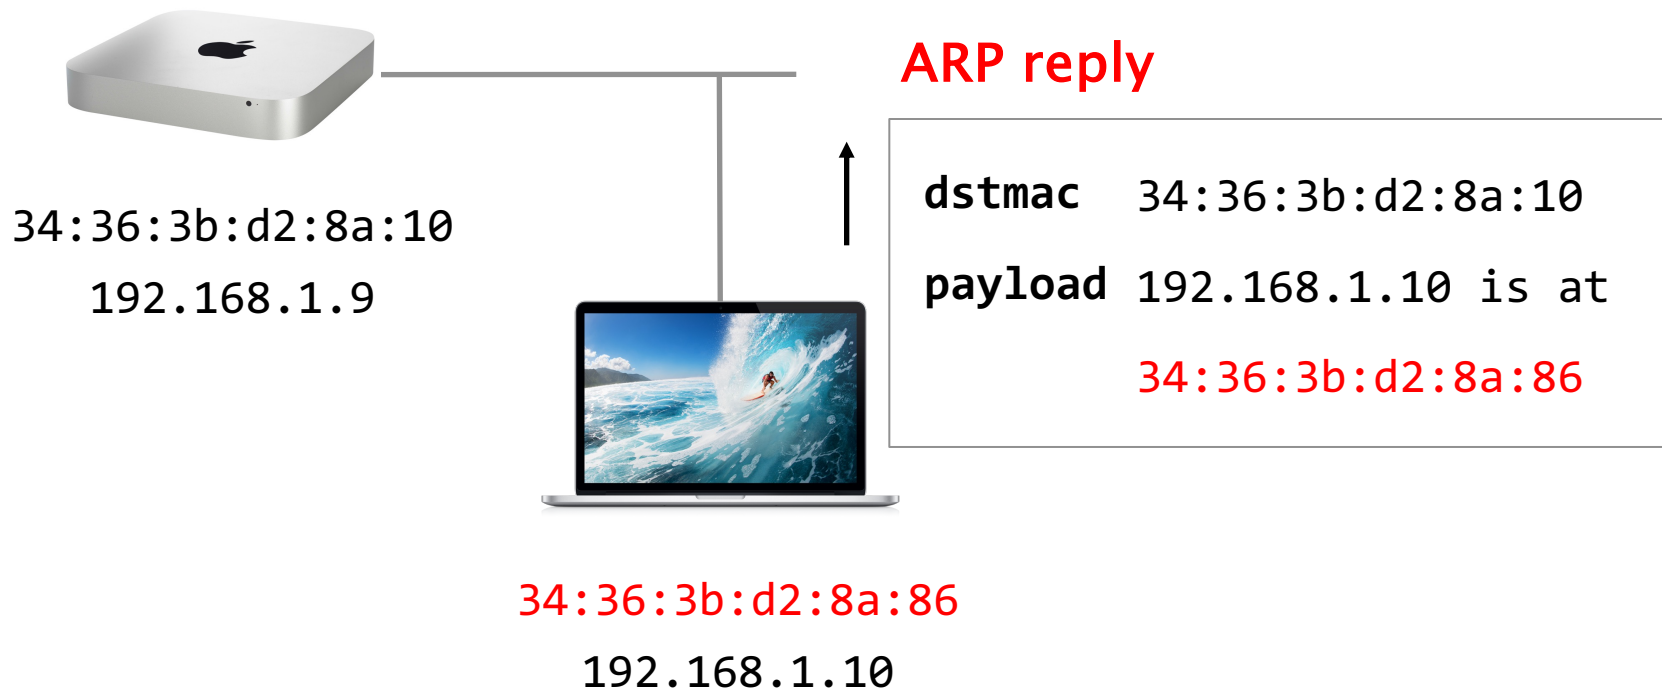
Address Resolution Protocol (ARP)

ARP request

```
dstmac  ff:ff:ff:ff:ff:ff
payload Who has 192.168.1.10?
        Tell 192.168.1.9
```




Address Resolution Protocol (ARP)



Address Resolution Protocol (ARP)

ARP table


192.168.1.10	34:36:3b:d2:8a:86
...	...



Ethernet Hub

Inefficient: each bit is sent everywhere

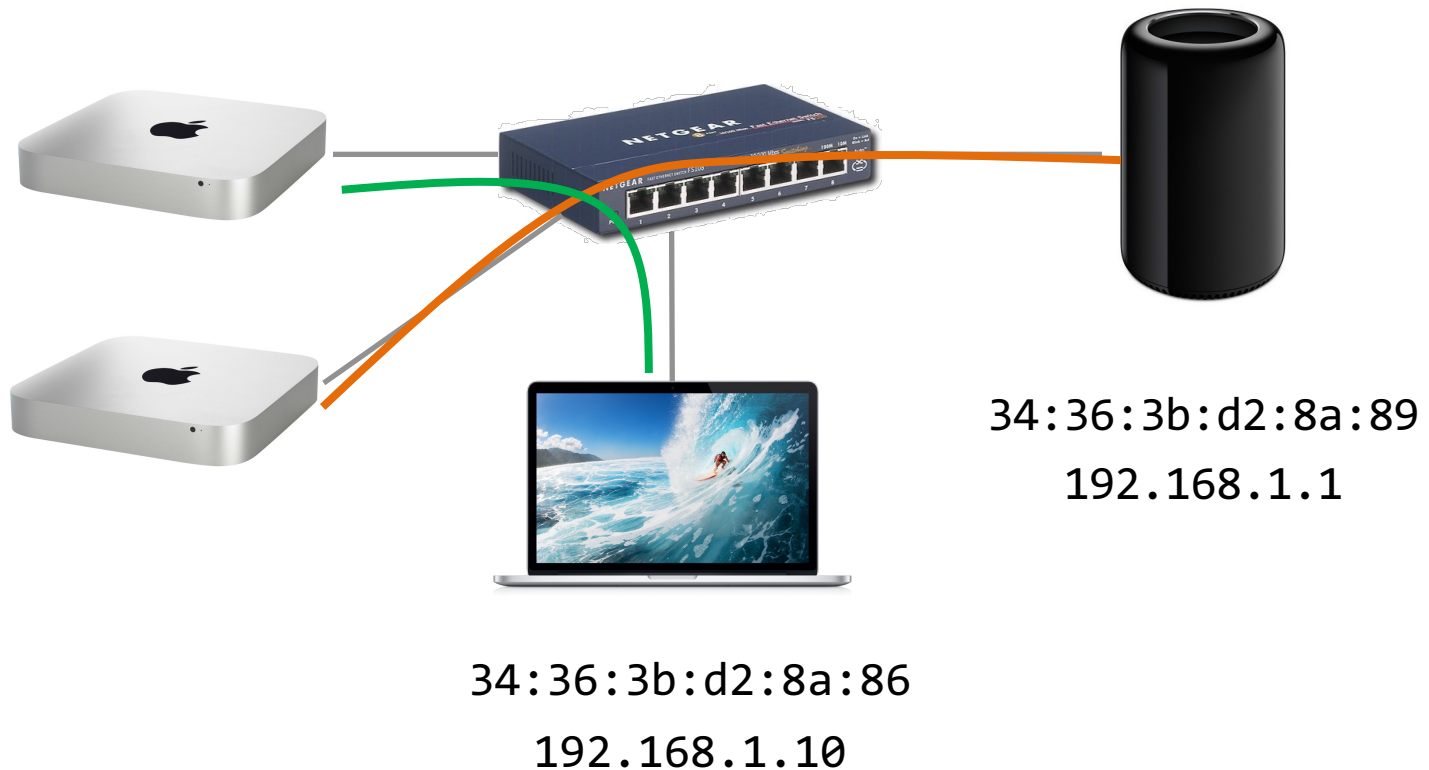
~~Obsolete~~
Ethernet hub
repeating bits from one port to
all the other ones



Ethernet Switch

efficient: communications
between different hosts
are isolated


Ethernet switch
repeating bits from one port to
the correct port



Ethernet Switch

- Switch is a plug-and-play device （即插即用设备）

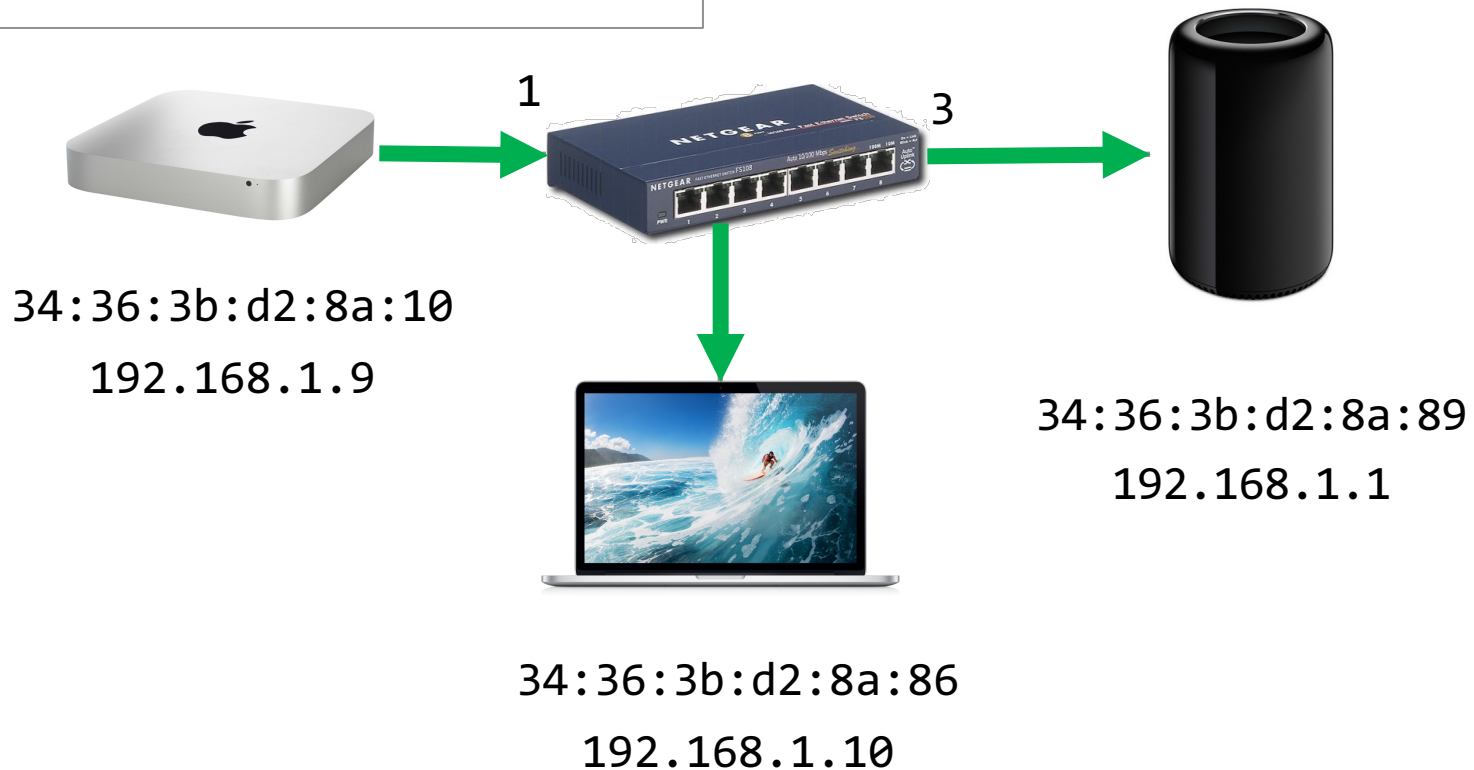
How does the switch know the correct port?



MAC Learning

ARP request

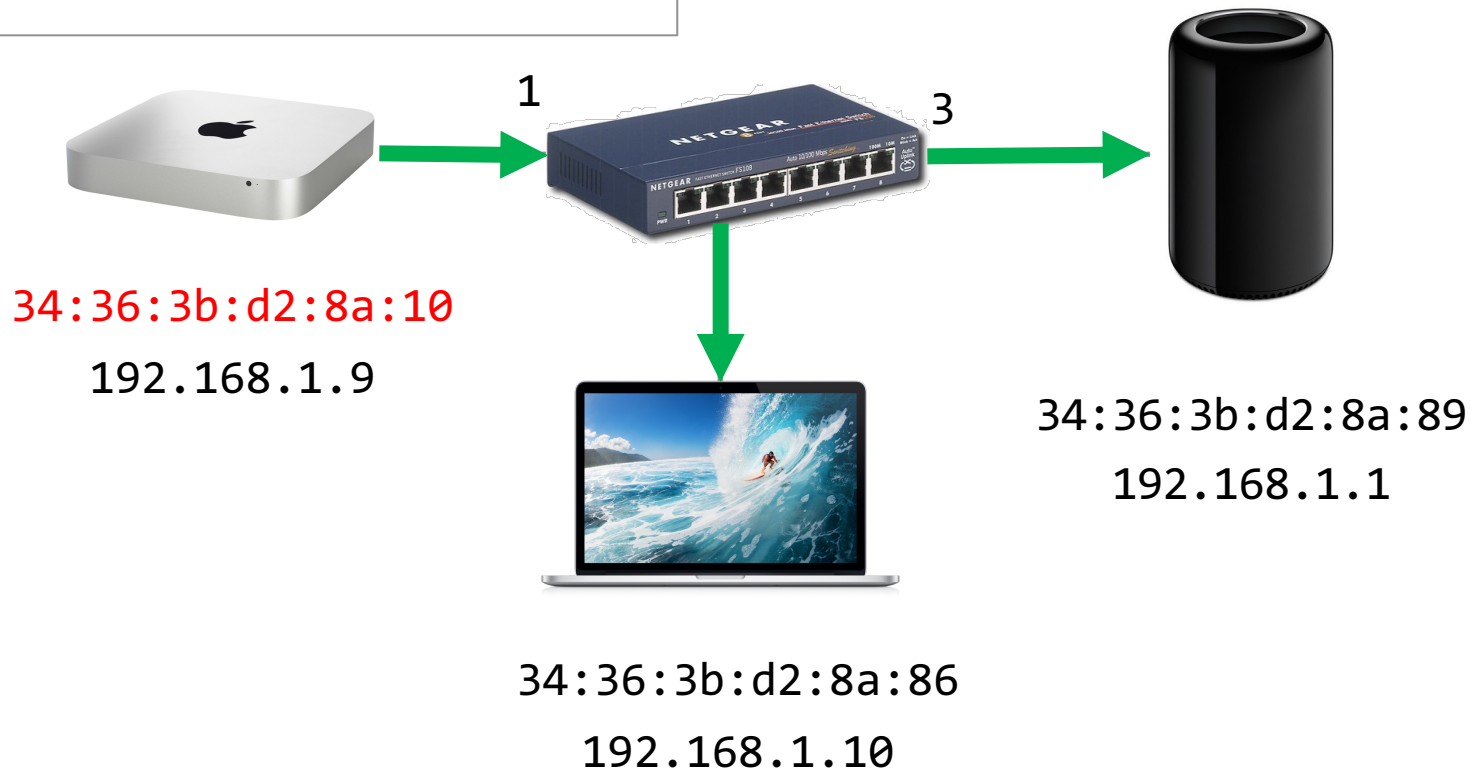
```
dstmac  ff:ff:ff:ff:ff:ff
payload Who has 192.168.1.10?
        Tell 192.168.1.9
```



MAC Learning


MAC Table

Dstmac	port
34:36:3b:d2:8a:10	1

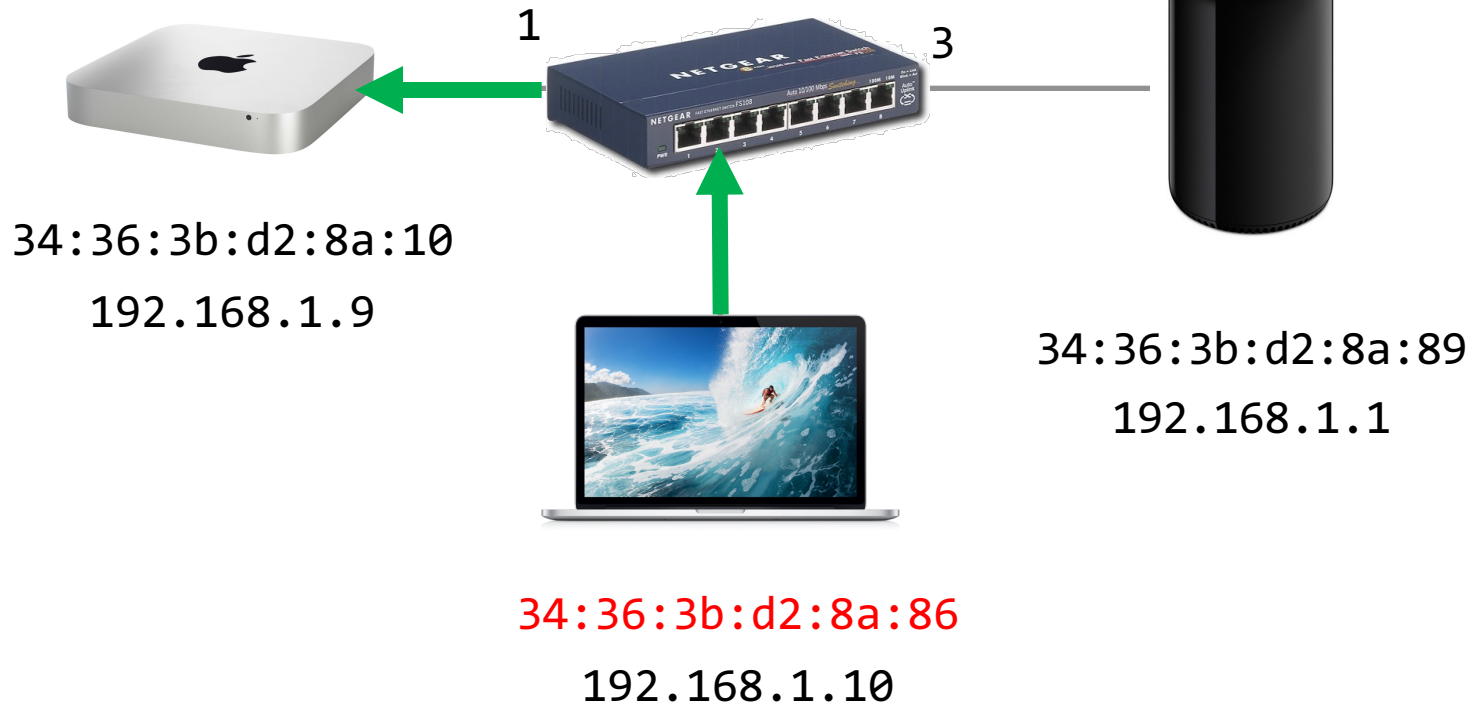


MAC Learning

MAC Table



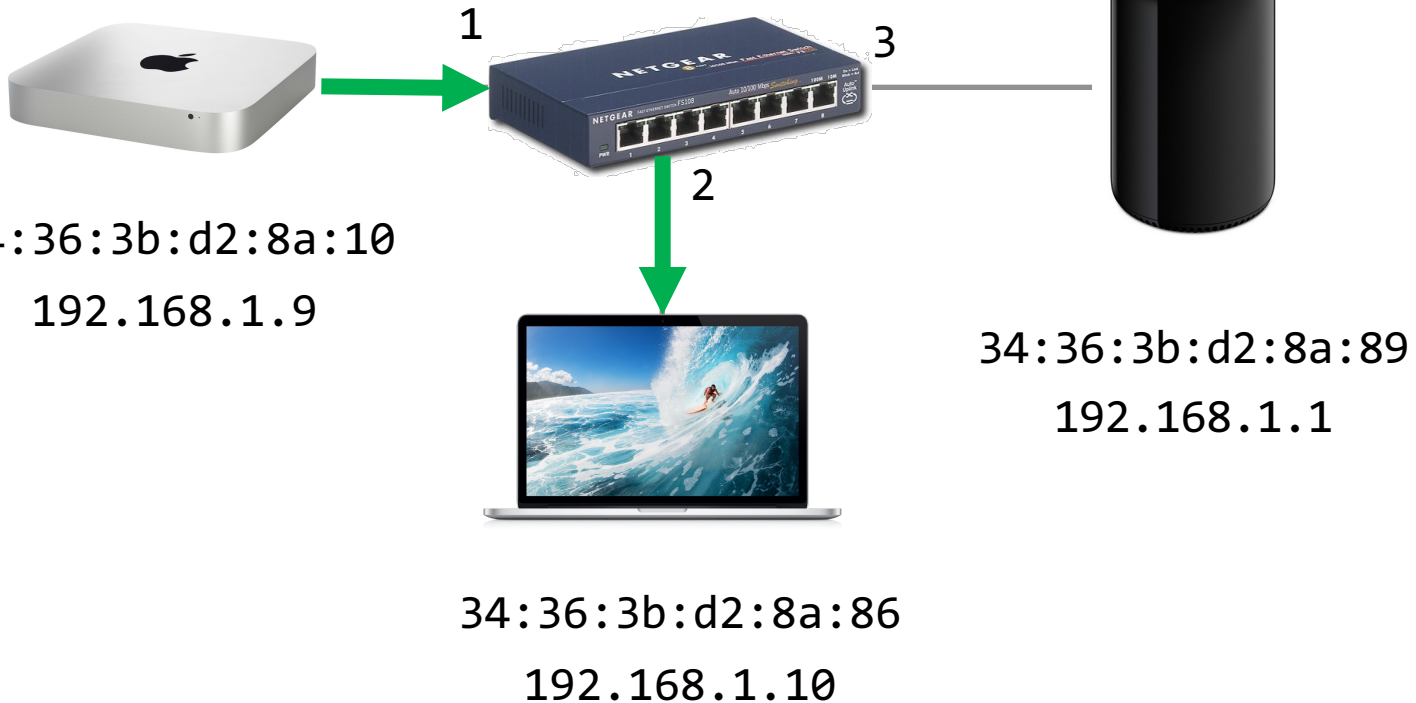
Dstmac	port
34:36:3b:d2:8a:10	1
34:36:3b:d2:8a:86	2




MAC Learning

MAC Table

Dstmac	port
34:36:3b:d2:8a:10	1
34:36:3b:d2:8a:86	2




Broadcast Storm



Each frame leads to the creation of *at least two new frames!*
exponential increase, with no TTL to remove looping frames...

Spanning Tree Protocol (STP)




Run the spanning tree protocol to disable some ports, so that the topology becomes a tree (rather than a graph)

What you actually do


How can I know www.example.com
is at 192.168.1.10?

I want to visit
www.example.com




Domain Name System (DNS)

- The DNS system is a distributed database which enables to resolve a **name** into an **IP address**

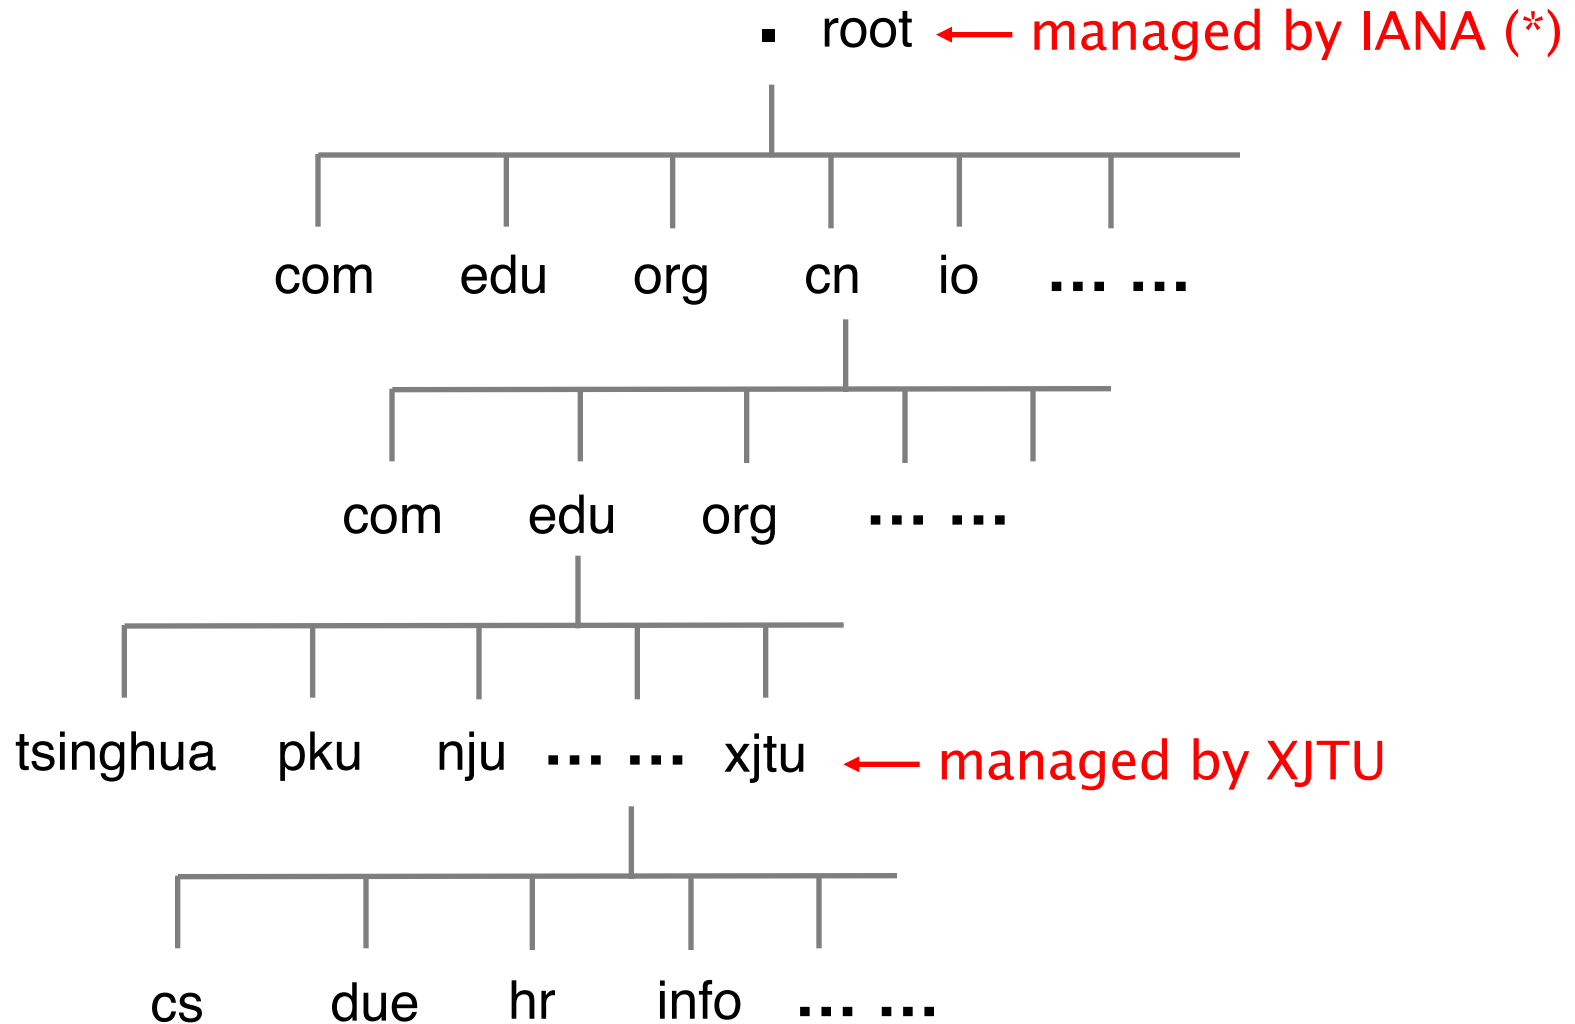


Domain Name System (DNS)



The **Domain Name System (DNS)** protocol is an important part of the web's infrastructure, serving as the **Internet's phone book**: every time you visit a website, your computer performs a DNS lookup.

Domain Name is Hierarchical



The Root

13 root
servers
across the
world

a. root-servers.net	VeriSign, Inc.
b. root-servers.net	University of Southern California
c. root-servers.net	Cogent Communications
d. root-servers.net	University of Maryland
e. root-servers.net	NASA
f. root-servers.net	Internet Systems Consortium
g. root-servers.net	US Department of Defense
h. root-servers.net	US Army
i. root-servers.net	Netnod
j. root-servers.net	VeriSign, Inc.
k. root-servers.net	RIPE NCC
l. root-servers.net	ICANN
m. root-servers.net	WIDE Project

DNS Server

- A DNS server stores Resource Records composed of a (name, value, type, TTL)

Records	Name	Value
A	hostname	IP address
NS	domain	DNS server name
MX	domain	Mail server name
CNAME	alias	canonical name
PTR	IP address	corresponding hostname

Using DIG to Resolve

- *dig* (domain information groper)

```
test@sdnexp:~$ dig
; <<>> DiG 9.11.3-lubuntu1.15-Ubuntu <<>>
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 40301
;; flags: qr rd ra; QUERY: 1, ANSWER: 13, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 65494
;; QUESTION SECTION:
; .                        IN      NS

;; ANSWER SECTION:
.      337324 IN      NS      j.root-servers.net.
.      337324 IN      NS      d.root-servers.net.
.      337324 IN      NS      e.root-servers.net.
.      337324 IN      NS      f.root-servers.net.
.      337324 IN      NS      h.root-servers.net.
.      337324 IN      NS      l.root-servers.net.
.      337324 IN      NS      g.root-servers.net.
.      337324 IN      NS      c.root-servers.net.
.      337324 IN      NS      m.root-servers.net.
.      337324 IN      NS      k.root-servers.net.
.      337324 IN      NS      b.root-servers.net.
.      337324 IN      NS      a.root-servers.net.
.      337324 IN      NS      i.root-servers.net.

;; Query time: 2 msec
;; SERVER: 127.0.0.53#53(127.0.0.53)
;; WHEN: Tue Feb 28 15:30:00 CST 2023
;; MSG SIZE rcvd: 239
```

```
test@sdnexp:~$ dig www.baidu.com
```

```
; <<>> DiG 9.11.3-1ubuntu1.15-Ubuntu <<>> www.baidu.com
```

```
;; global options: +cmd
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 44151
```

```
;; flags: qr rd ra; QUERY: 1, ANSWER: 3, AUTHORITY: 0, ADDITIONAL: 1
```

```
;; OPT PSEUDOSECTION:
```

```
; EDNS: version: 0, flags:; udp: 65494
```

```
;; QUESTION SECTION:
```

```
;www.baidu.com.
```

IN

A

Internet

规范名字

```
;; ANSWER SECTION:
```

```
www.baidu.com.
```

699

IN

CNAME

www.a.shifen.com.

```
www.a.shifen.com.
```

235

IN

A

182.61.200.7

```
www.a.shifen.com.
```

235

IN

A

182.61.200.6

```
;; Query time: 2 msec
```

```
;; SERVER: 127.0.0.53#53(127.0.0.53)
```

```
;; WHEN: Tue Feb 28 15:32:40 CST 2023
```

```
;; MSG SIZE rcvd: 101
```

What you actually do

DNS request

```
dstip 192.168.1.1
payload What is the IP address
of www.example.com?
```

DNS Server
UDP Port 53



34:36:3b:d2:8a:10
192.168.1.9



34:36:3b:d2:8a:89
192.168.1.1



34:36:3b:d2:8a:86
192.168.1.10
www.example.com

What you actually do

DNS Reply

`dstip` 192.168.1.9
`payload` The IP for www.example.com
is 192.168.1.10

DNS Server
UDP Port 53



34:36:3b:d2:8a:10
192.168.1.9



34:36:3b:d2:8a:89
192.168.1.1



34:36:3b:d2:8a:86
192.168.1.10
www.example.com

Public DNS Servers

DNS request

dstip 8.8.8.8

payload What is the IP address
of www.baidu.com?

Google Public DNS



34:36:3b:d2:8a:10
192.168.1.9



34:36:3b:d2:8a:86
192.168.1.10
www.example.com



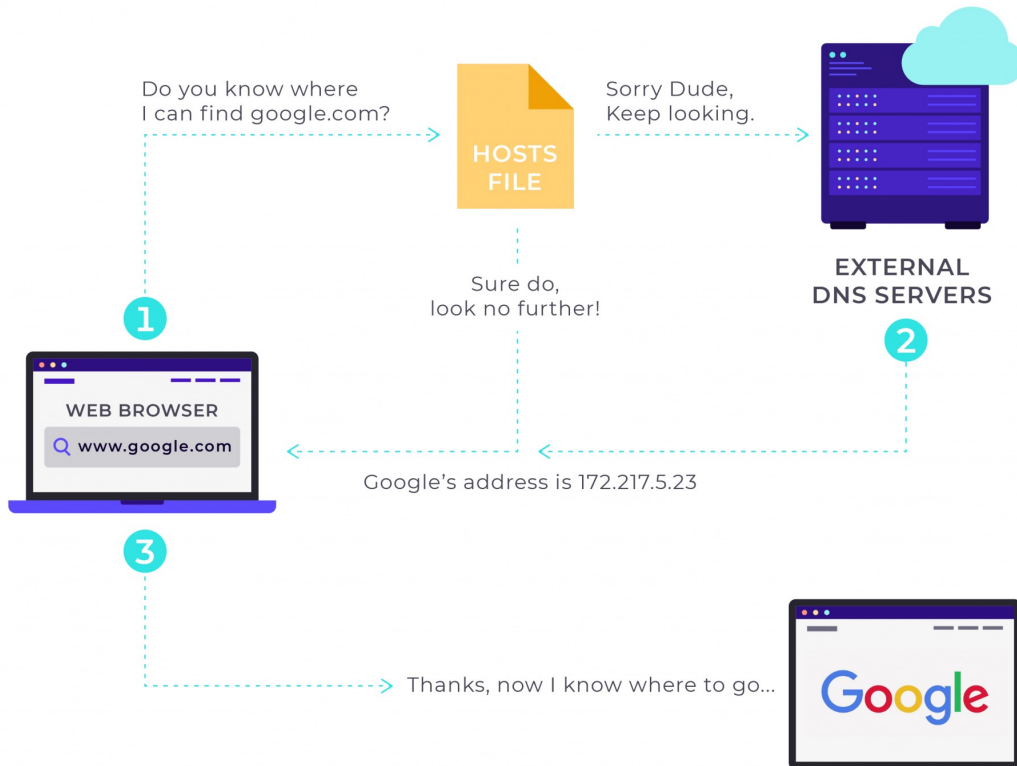
34:36:3b:d2:8a:89
192.168.1.1

Public DNS Servers

- Cloudflare: 1.1.1.1
- 阿里云: 223.5.5.5
- 百度: 180.76.76.76
- 腾讯: 119.29.29.29

The hosts File

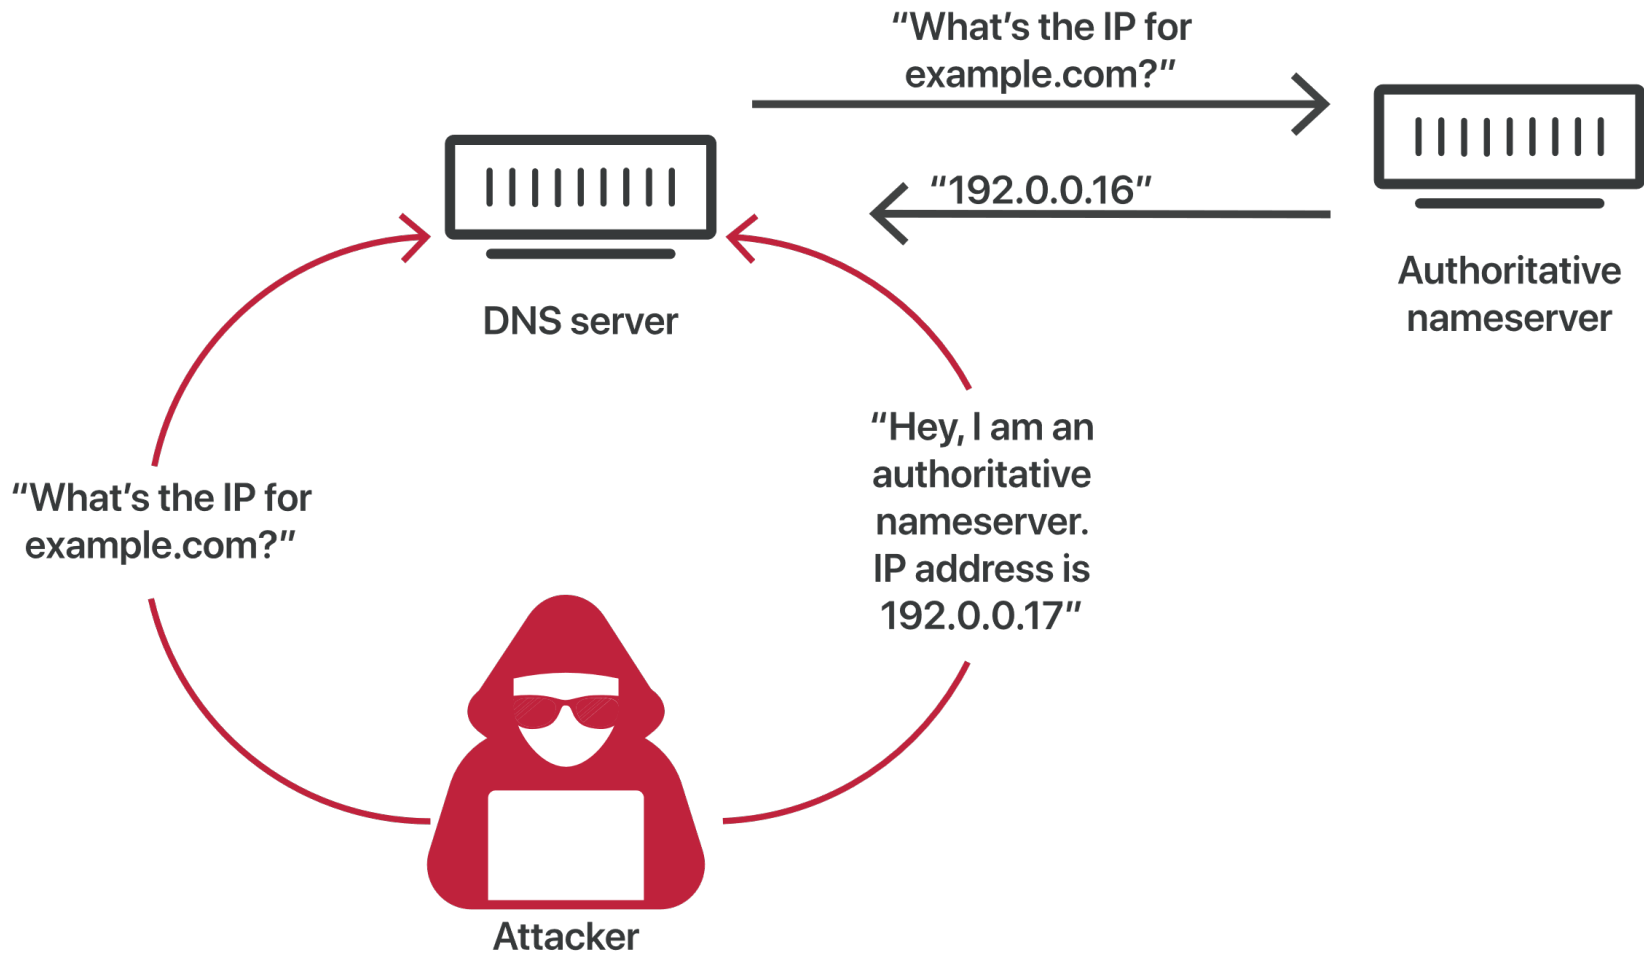
How a hosts file works



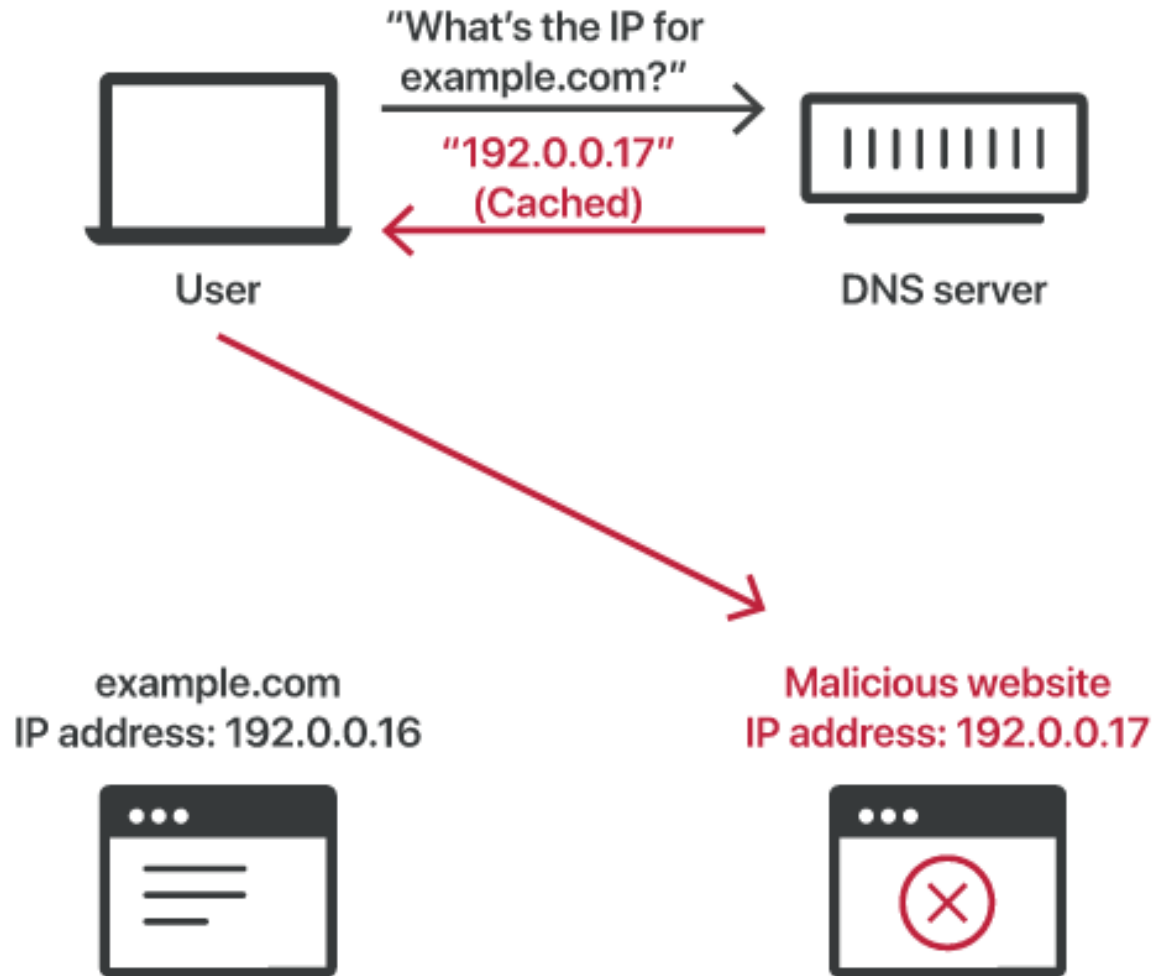
Host file takes priority over DNS

- How to modify it?
 - Linux/MacOS: /etc/hosts
 - Windows: c:\Windows\System32\Drivers\etc\hosts

DNS Cache Poisoning



DNS Cache Poisoning



DNS Cache Poisoning



How to Reach the Outside World

DNS request

```
dstip 8.8.8.8
```

```
payload What is the IP address  
of www.baidu.com?
```



34:36:3b:d2:8a:10

192.168.1.9

Subnet: 255.255.255.0

8.8.8.8 is not
in the network



34:36:3b:d2:8a:86

192.168.1.10

www.example.com



34:36:3b:d2:8a:89

192.168.1.1

How to Reach the Outside World

DNS request

dstmac 34:36:3b:d2:8a:89

dstip 8.8.8.8

payload What is the IP address
of www.baidu.com?



34:36:3b:d2:8a:10

192.168.1.9

Subnet: 255.255.255.0

Default gateway: 192.168.1.1



34:36:3b:d2:8a:89

192.168.1.1



34:36:3b:d2:8a:86

192.168.1.10

www.example.com

How to Reach the Outside World

