

SDN-Lab2



本实验工作目录：

```
> pwd
/home/root-hbx/CS

> ls
__pycache__  Broadcast_Loop.py  gmon.out  Learning_Switch.py  ryu-xjt
```

1. 自学习交换机

1.1 问题说明



情景：

1969年的 ARPANET 非常简单，仅由四个结点组成。假设每个结点都对应一个交换机，每个交换机都具有 一个直连主机，你的任务是实现不同主机之间的正常通信。

前文给出的简单交换机洪泛数据包，虽然能初步实现主机间的通信，但会带来不必要的带宽消耗，并且 会使通信内容泄露给第三者。因此，请你在简单交换机的基础上实现二层自学习交换机，避免数据包的 洪泛。

SDN 自学习交换机的工作流程可以参考：

1. 控制器为每个交换机维护一个 mac - port 映射表
2. 控制器收到 packet_in 消息后，解析其中携带的数据包
3. 控制器学习 src_mac - in_port 映射
4. 控制器查询 dst_mac ，如果未学习，则洪泛数据包；如果已学习，则向指定端口转发数据包 (packet_out)，并向交换机下发流表项(flow_mod)，指导交换机转发同类型的数据包。

采用 ryu 作为远程控制器 (remote controller)，开启方式：

```
ryu-manager --observe-links Broadcast_Loop.py
```

网络拓扑为 topo_1969_1.py ，启动方式：

```
sudo mn --custom topo_1969_1.py --topo generated --controller remote
```

1.2 实验代码

task1_selfLearningSwitch.py:

```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import MAIN_DISPATCHER, CONFIG_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
from ryu.lib.packet import arp
from ryu.lib.packet import ether_types

ETHERNET = ethernet.ethernet.__name__
ETHERNET_MULTICAST = "ff:ff:ff:ff:ff:ff"
ARP = arp.arp.__name__
```

```

'''
an adaptation in Broadcast_Loop.py
'''

class Switch_Dict(app_manager.RyuApp):
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]

    def __init__(self, *args, **kwargs):
        super(Switch_Dict, self).__init__(*args, **kwargs)
        self.sw = {}
        # topo: src--in_port--|Sw = dpid|--aim_port--dst
        # mapping-1: portIn = mac_to_port[Sw][src]
        # mapping-2: portOut = mac_to_port[Sw][dst]

        self.mac_to_port = {}

    def add_flow(self, datapath, priority, match, actions, idle_timeout=0, ha
        dp = datapath
        ofp = dp.ofproto
        parser = dp.ofproto_parser
        inst = [parser.OFPInstructionActions(ofp.OFPIT_APPLY_ACTIONS, actions
        mod = parser.OFPFlowMod(datapath=dp, priority=priority,
                                idle_timeout=idle_timeout,
                                hard_timeout=hard_timeout,
                                match=match, instructions=inst)

        dp.send_msg(mod)

    @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
    def switch_features_handler(self, ev):
        msg = ev.msg
        dp = msg.datapath
        ofp = dp.ofproto
        parser = dp.ofproto_parser
        match = parser.OFPMatch()
        actions = [parser.OFPActionOutput(ofp.OFPP_CONTROLLER, ofp.OFPCML_NO_
        self.add_flow(dp, 0, match, actions)

    # what we actually done:
    @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
    def packet_in_handler(self, ev):
        # method-def (Packet-in): self -- An event

        msg = ev.msg # message object
        dp = msg.datapath # data object
        ofp = dp.ofproto # constants about OpenFlow

```

```

parser = dp.ofproto_parser # parser to construct and analysis OpenFlo

# the identity of switch
dpid = dp.id # Sw.id
self.mac_to_port.setdefault(dpid, {}) # add {empty} into Sw.

# the port that receive the packet
in_port = msg.match['in_port']
pkt = packet.Packet(msg.data)
eth_pkt = pkt.get_protocol(ethernet.ethernet)
if eth_pkt.ethertype == ether_types.ETH_TYPE_LLDP:
    return
if eth_pkt.ethertype == ether_types.ETH_TYPE_IPV6:
    return

# get the mac
dst = eth_pkt.dst
src = eth_pkt.src

# get protocols
header_list = dict((p.protocol_name, p) for p in pkt.protocols if typ
if dst == ETHERNET_MULTICAST and ARP in header_list:
    # need to code here to avoid broadcast loop to finish mission 2
    # this part can be passed here
    pass

# self-learning
# src---in_port---|Sw = dpid|---aim_port---dst

self.mac_to_port[dpid][src] = in_port # mapping-1
currentSw = self.mac_to_port[dpid]
flag = 0 # judge if it's a new mapping

# the logic process of OpenFlow Controller

if dst in currentSw:
    aim_port = self.mac_to_port[dpid][dst] # mapping-2
else:
    aim_port = ofp.OFPP_FLOOD # pre-flooding this packet to all nodes
    flag = 1

# sending to which port (specific one / flooding) depends on "aim_por
actions = [parser.OFPAActionOutput(aim_port)]

# a new mapping, add the flow table to switch
if (flag):

```

```

        # create a matching condition, only the pkt (portIn: in_port, dst
        portMacPair = parser.OFPMatch(in_port=in_port, eth_dst=dst)
        # add this matching into flow table
        self.add_flow(dp, 10, portMacPair, actions) # avoid flow table sh

'''
After receiving Pack-In,
controller will send Packet-Out to Switch and give orders to the spec
'''

data = None
# for flow tables, in-time send is necessary
if msg.buffer_id == ofp.OFP_NO_BUFFER:
    # if there's no packetBuffer in Switch, it implies:
    #   this packet is sending to Controller directly!
    #   and we need to copy the "data in packet" into dataMessage
    #   for it will be loaded into "outPut" towards Sw.
    data = msg.data

outPut = parser.OFPPacketOut(
    datapath=dp, # towards the Sw. who is sending Packet-In message
    buffer_id=msg.buffer_id, # the Buffer ID of the packetNum
    in_port=in_port, # the receiving portNum
    actions=actions,
    data=data,
)

dp.send_msg(outPut)

```

这段代码有一个隐藏的问题，不会影响问题1，但是会严重影响问题2

```

# the logic process of OpenFlow Controller

if dst in currentSw:
    aim_port = self.mac_to_port[dpid][dst] # mapping-2
else:
    aim_port = ofp.OFPP_FLOOD # pre-flooding this packet to all nodes
    flag = 1

# sending to which port (specific one / flooding) depends on "aim_por
actions = [parser.OFPACTIONOutput(aim_port)]

# a new mapping, add the flow table to switch
if (flag):
    # create a matching condition, only the pkt (portIn: in_port, dst

```

```
portMacPair = parser.OFPMatch(in_port=in_port, eth_dst=dst)
# add this matching into flow table
self.add_flow(dp, 10, portMacPair, actions) # avoid flow table sh
```

问题2做出了对应的修改

1.3 实验分析

原理

当前方案中给出的"自学习", 本质上是根据"映射关系"进行查询:

- 当交换机上报一个Packet In消息给控制器后, 控制器检查该消息携带的是否为Ethernet类型报文
 - 如果是: 提取出 eth_src 和 portIn, 建立映射关系
 - 反之: 不建立
- 后续Ethernet类型报文进入时, 控制器检测是否已学习过该报文中 dst_mac 对应的portIn
 - 如果是: 建立对应关系, 并下发FlowTable
 - 反之: 说明当前是ARP, 立即洪泛

安装控制器并进行对接

```
ryu-manager --observe-links task1_selfLearningSwitch.py
```

^ ~/CS

^ ~/CS

```
> ryu-manager --observe-links task1_selfLearningSwitch.py
loading app task1_selfLearningSwitch.py
loading app ryu.controller.ofp_handler
instantiating app task1_selfLearningSwitch.py of Switch_Dict
instantiating app ryu.controller.ofp_handler of OFPHandler
█
```

SDN_Spring_2
SDN_Spring_2024 17:32:49

pingall 结果

```
sudo mn --custom topo_1969_1.py --topo generated --controller remote
```

```
mininet> pingall
```

```
mininet> pingall
*** Ping: testing ping reachability
SR1 -> *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节, 来自 10.0.0.2: icmp_seq=1 ttl=64 时间=6.56 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 6.563/6.563/6.563/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节, 来自 10.0.0.3: icmp_seq=1 ttl=64 时间=6.35 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 6.345/6.345/6.345/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节, 来自 10.0.0.4: icmp_seq=1 ttl=64 时间=4.04 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 4.037/4.037/4.037/0.000 ms

X
UCLA -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节, 来自 10.0.0.1: icmp_seq=1 ttl=64 时间=2.18 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 2.177/2.177/2.177/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节, 来自 10.0.0.3: icmp_seq=1 ttl=64 时间=3.57 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 3.572/3.572/3.572/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节, 来自 10.0.0.4: icmp_seq=1 ttl=64 时间=3.65 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 3.654/3.654/3.654/0.000 ms

X
UCSB -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节, 来自 10.0.0.1: icmp_seq=1 ttl=64 时间=2.29 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 2.292/2.292/2.292/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节, 来自 10.0.0.2: icmp_seq=1 ttl=64 时间=3.24 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 3.236/3.236/3.236/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节, 来自 10.0.0.4: icmp_seq=1 ttl=64 时间=2.78 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 2.780/2.780/2.780/0.000 ms

X
UTAH -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节, 来自 10.0.0.1: icmp_seq=1 ttl=64 时间=3.06 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 3.063/3.063/3.063/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节, 来自 10.0.0.2: icmp_seq=1 ttl=64 时间=2.75 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 2.752/2.752/2.752/0.000 ms

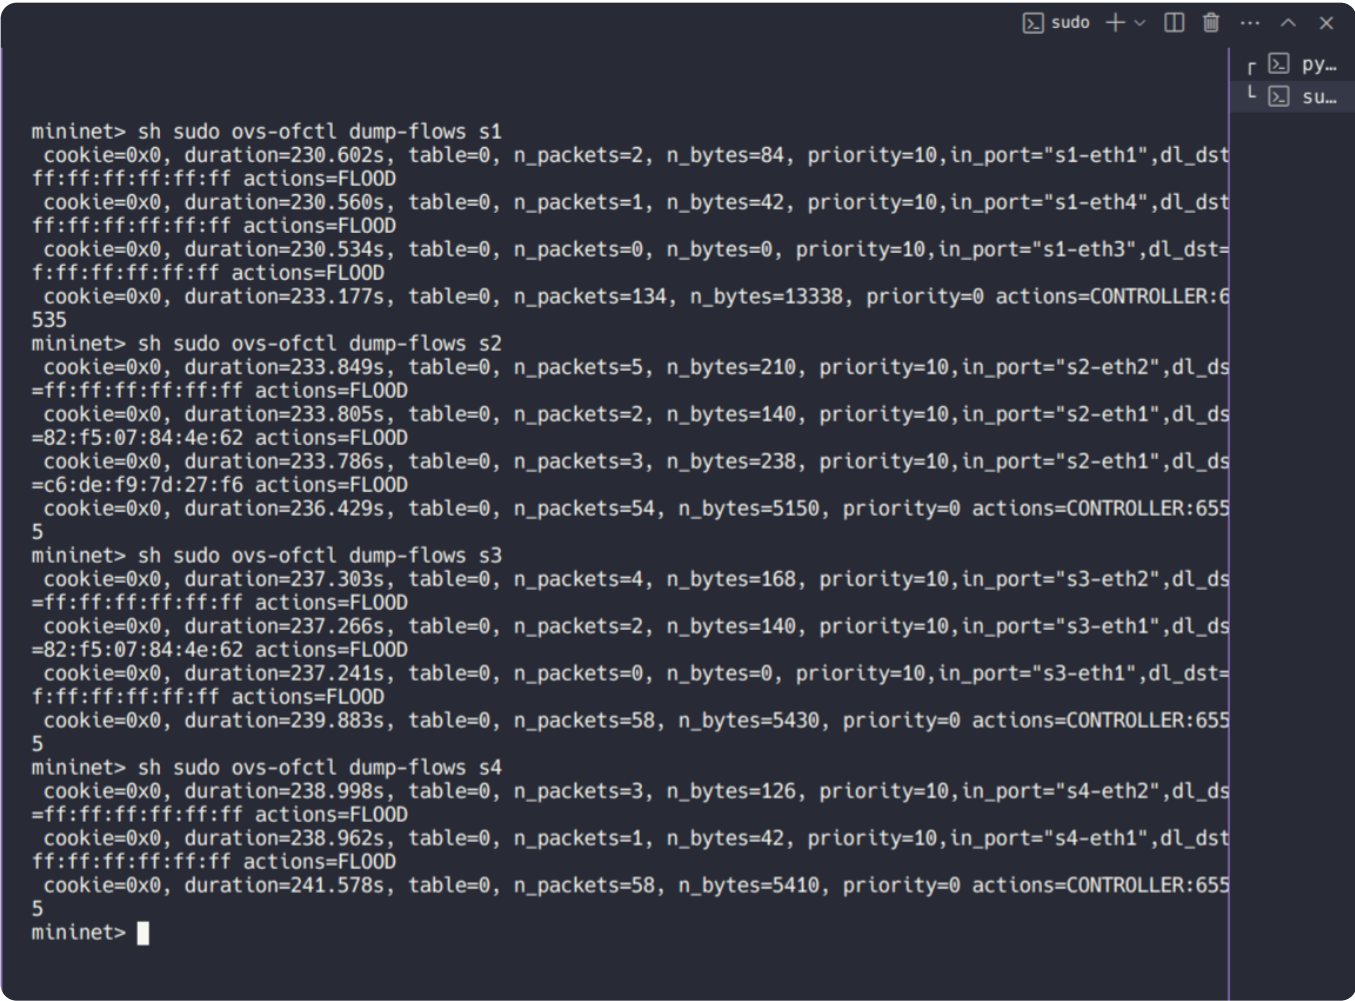
X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节, 来自 10.0.0.3: icmp_seq=1 ttl=64 时间=2.70 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 2.698/2.698/2.698/0.000 ms

X
*** Results: 100% dropped (0/12 received)
mininet>
Interrupt
```

查看各Switch流表



```
sudo ovs-ofctl dump-flows s1 # SRI
sudo ovs-ofctl dump-flows s2 # UCLA
sudo ovs-ofctl dump-flows s3 # UCSB
sudo ovs-ofctl dump-flows s4 # UTAH
```



```
mininet> sh sudo ovs-ofctl dump-flows s1
 cookie=0x0, duration=230.602s, table=0, n_packets=2, n_bytes=84, priority=10,in_port="s1-eth1",dl_dst=
 ff:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=230.560s, table=0, n_packets=1, n_bytes=42, priority=10,in_port="s1-eth4",dl_dst=
 ff:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=230.534s, table=0, n_packets=0, n_bytes=0, priority=10,in_port="s1-eth3",dl_dst=
 f:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=233.177s, table=0, n_packets=134, n_bytes=13338, priority=0 actions=CONTROLLER:6
 535
mininet> sh sudo ovs-ofctl dump-flows s2
 cookie=0x0, duration=233.849s, table=0, n_packets=5, n_bytes=210, priority=10,in_port="s2-eth2",dl_ds
 =ff:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=233.805s, table=0, n_packets=2, n_bytes=140, priority=10,in_port="s2-eth1",dl_ds
 =82:f5:07:84:4e:62 actions=FL00D
 cookie=0x0, duration=233.786s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s2-eth1",dl_ds
 =c6:de:f9:7d:27:f6 actions=FL00D
 cookie=0x0, duration=236.429s, table=0, n_packets=54, n_bytes=5150, priority=0 actions=CONTROLLER:655
 5
mininet> sh sudo ovs-ofctl dump-flows s3
 cookie=0x0, duration=237.303s, table=0, n_packets=4, n_bytes=168, priority=10,in_port="s3-eth2",dl_ds
 =ff:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=237.266s, table=0, n_packets=2, n_bytes=140, priority=10,in_port="s3-eth1",dl_ds
 =82:f5:07:84:4e:62 actions=FL00D
 cookie=0x0, duration=237.241s, table=0, n_packets=0, n_bytes=0, priority=10,in_port="s3-eth1",dl_dst=
 f:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=239.883s, table=0, n_packets=58, n_bytes=5430, priority=0 actions=CONTROLLER:655
 5
mininet> sh sudo ovs-ofctl dump-flows s4
 cookie=0x0, duration=238.998s, table=0, n_packets=3, n_bytes=126, priority=10,in_port="s4-eth2",dl_ds
 =ff:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=238.962s, table=0, n_packets=1, n_bytes=42, priority=10,in_port="s4-eth1",dl_dst
 ff:ff:ff:ff:ff:ff actions=FL00D
 cookie=0x0, duration=241.578s, table=0, n_packets=58, n_bytes=5410, priority=0 actions=CONTROLLER:655
 5
mininet> █
```

2. 避免环路广播

2.1 问题说明



情景：

UCLA 和 UCSB 通信频繁，两者间建立了一条直连链路。

在新的拓扑 `topo_1969_2.py` 中运行自学习交换机，UCLA 和 UTAH 之间无法正常通信。分析流表发现，源主机虽然只发了很少的几个数据包，但流表项却匹配了上千次；WireShark 也截取到了数目异常大的相同报文

这实际上是 ARP 广播数据包在环状拓扑中洪泛导致的，传统网络利用生成树协议解决这一问题。

在 SDN 中，不必局限于生成树协议，可以通过多种新的策略解决这一问题。以下给出一种解决思路，请在自学习交换机的基础上完善代码，解决问题：

1. 当序号为 `dpid` 的交换机从 `portIn` 第一次收到某个 `src_mac` 主机发出，询问 `dst_ip` 的广播 ARP Request 数据包时，控制器记录一个映射 `(dpid, src_mac, dst_ip) -> in_port`；
2. 下一次该交换机收到同一 `(src_mac, dst_ip)` 但 `in_port` 不同的 ARP Request 数据包时直接丢弃即可

2.2 实验代码

task2_BreakLoop.py:

PYTHON

```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import MAIN_DISPATCHER, CONFIG_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
from ryu.lib.packet import arp
from ryu.lib.packet import ether_types

ETHERNET = ethernet.ethernet.__name__
ETHERNET_MULTICAST = "ff:ff:ff:ff:ff:ff"
ARP = arp.arp.__name__

class Switch_Dict(app_manager.RyuApp):
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]

    def __init__(self, *args, **kwargs):
        super(Switch_Dict, self).__init__(*args, **kwargs)

        self.arp_in_port = {}
        self.mac_to_port = {}

    def add_flow(
        self, datapath, priority, match, actions, idle_timeout=0, hard_timeou
    ):
        dp = datapath
        ofp = dp.ofproto
        parser = dp.ofproto_parser

        inst = [parser.OFPInstructionActions(ofp.OFPIT_APPLY_ACTIONS, actions
        mod = parser.OFPFlowMod(
            datapath=dp,
            priority=priority,
            idle_timeout=idle_timeout,
            hard_timeout=hard_timeout,
            match=match,
            instructions=inst,
        )
        dp.send_msg(mod)

    @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
    def switch_features_handler(self, ev):
```

```

        msg = ev.msg
        dp = msg.datapath
        ofp = dp.ofproto
        parser = dp.ofproto_parser
        match = parser.OFPMatch()
        actions = [parser.OFPActionOutput(ofp.OFPP_CONTROLLER, ofp.OFPCML_NO_
self.add_flow(dp, 0, match, actions)

# what we actually done:
@set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
def packet_in_handler(self, ev):
    msg = ev.msg
    dp = msg.datapath
    ofp = dp.ofproto
    parser = dp.ofproto_parser

    # the identity of switch
    dpid = dp.id
    self.mac_to_port.setdefault(dpid, {})
    self.arp_in_port.setdefault(dpid, {})

    # the port that receive the packet
    in_port = msg.match["in_port"]
    pkt = packet.Packet(msg.data)
    eth_pkt = pkt.get_protocol(ethernet.ethernet)
    if eth_pkt.ethertype == ether_types.ETH_TYPE_LLDP:
        return
    if eth_pkt.ethertype == ether_types.ETH_TYPE_IPV6:
        return

    # get the mac
    eth_dst = eth_pkt.dst
    eth_src = eth_pkt.src

    # get protocols
    header_list = dict(
        (p.protocol_name, p) for p in pkt.protocols if type(p) != str
    )

    # ARP Loop Processing
    if eth_dst == ETHERNET_MULTICAST and ARP in header_list:
        # a ARP packet
        arp_pkt = pkt.get_protocol(arp.arp)

        # ARP request packet
        if arp_pkt and arp_pkt.opcode == arp.ARP_REQUEST:

```

```

# arp_pkt.opcode represents the option-code of ARP_Packet
# option-code: the action message need to transfer back to Sw
# arp.ARP_REQUEST: Sw. is waiting for {"the MAC Addr." which
req_dst_ip = arp_pkt.dst_ip
arp_src_mac = arp_pkt.src_mac

...

ARP: need the dstMAC according to dstIP

(srcMAC, dstIP, Sw.pid)
- if srcMAC is same && dstIP is same
    - if portIn is same: pass
    - else: "what we have used" => discard
...

# learned the mac in mapping
if arp_src_mac in self.arp_in_port[dpid]:
    # if the srcMAC is recorded

    # get IP in mapping
    if req_dst_ip in self.arp_in_port[dpid][arp_src_mac]:
        # if the dstIP is also recorded

        if in_port != self.arp_in_port[dpid][arp_src_mac][req
            match = parser.OFPMatch(
                in_port=in_port,
                arp_op=arp.ARP_REQUEST,
                arp_tpa=req_dst_ip,
                arp_sha=arp_src_mac,
            )
            actions = []
            # higher than self-learning
            self.add_flow(dp, 20, match, actions)

            out = parser.OFPPacketOut(
                datapath=dp,
                buffer_id=msg.buffer_id,
                in_port=in_port,
                actions=[],
                data=None,
            )
            dp.send_msg(out)
            return

# no req_dst_ip in mapping
else:

```

```

        # record the dstIP and mapping(srcMAC, dstIP, portIn)
        self.arp_in_port[dpid][arp_src_mac].setdefault(req_dst_ip, in_port)
        self.arp_in_port[dpid][arp_src_mac][req_dst_ip] = in_port
    # no arp_src_mac in mapping
    else:
        # record the srcMAC and mapping(srcMAC, dstIP, portIn)
        self.arp_in_port[dpid].setdefault(arp_src_mac, {})
        self.arp_in_port[dpid][arp_src_mac][req_dst_ip] = in_port

# self-learning
self.mac_to_port[dpid][eth_src] = in_port
if eth_dst in self.mac_to_port[dpid]:
    # have learned the mac-port mapping
    out_port = self.mac_to_port[dpid][eth_dst]
else:
    # no, just flood
    out_port = ofp.OFPP_FLOOD
# output the packet
actions = [parser.OFPActionOutput(out_port)]

# if learned a new mapping, add the flow table to switch
if out_port != ofp.OFPP_FLOOD:
    match = parser.OFPMatch(in_port=in_port, eth_dst=eth_dst)
    # set priority to 10 to avoid flow table shadowing
    self.add_flow(dp, 10, match, actions)

data = None
# for flow tables, in-time send is necessary
if msg.buffer_id == ofp.OFP_NO_BUFFER:
    data = msg.data
out = parser.OFPPacketOut(
    datapath=dp,
    buffer_id=msg.buffer_id,
    in_port=in_port,
    actions=actions,
    data=data,
)
dp.send_msg(out)

```

```

# self-learning
self.mac_to_port[dpid][eth_src] = in_port
if eth_dst in self.mac_to_port[dpid]:
    # have learned the mac-port mapping
    out_port = self.mac_to_port[dpid][eth_dst]
else:

```

```
# no, just flood
out_port = ofp.OFPP_FLOOD
# output the packet
actions = [parser.OFPACTIONOutput(out_port)]

# if learned a new mapping, add the flow table to switch
if out_port != ofp.OFPP_FLOOD:
    match = parser.OFPMATCH(in_port=in_port, eth_dst=eth_dst)
    # set priority to 10 to avoid flow table shadowing
    self.add_flow(dp, 10, match, actions)
```

问题2的改正如上

原因分析：flag = 1,

2.3 实验分析

原理

当前方案给出的本质上还是 "映射记录,重复舍弃"

当交换机上报一个Packet In消息给控制器后, 控制器检查最基本条件:

1. 该消息携带的是否为 Ethernet 类型报文
2. 该数据包是否携带 ARP 报头

如上述条件满足,则进入映射设计:

- 我们采取针对某一个交换机考察, 建立映射 (srcMAC, dstIP, portIn)
- 具体来说, 设计: (srcMAC, dstIP, currentSw) = portIn_Sw

针对ARP的专门处理:

- 如果 *srcMAC is same && dstIP is same*
 - 如果 portIn 不相同, 说明是经环路运作后的"重复包", 直接丢弃即可
 - 反之, pass即可

安装控制器并进行对接

```
ryu-manager --observe-links task2_BreakLoop.py
```

```
A ~/CS SDN_Spring_2024 18:15:44
> ryu-manager --observe-links task2_BreakLoop.py
loading app task2_BreakLoop.py
loading app ryu.controller.ofp_handler
instantiating app task2_BreakLoop.py of Switch_Dict
instantiating app ryu.controller.ofp_handler of OFPHandler
^C

A ~/CS 8m 37s SDN_Spring_2024 18:26:32
> ryu-manager --observe-links task2_BreakLoop.py
loading app task2_BreakLoop.py
loading app ryu.controller.ofp_handler
instantiating app task2_BreakLoop.py of Switch_Dict
instantiating app ryu.controller.ofp_handler of OFPHandler
^C

A ~/CS 1m 46s SDN_Spring_2024 18:28:31
> ryu-manager --observe-links task2_BreakLoop.py
loading app task2_BreakLoop.py
loading app ryu.controller.ofp_handler
instantiating app task2_BreakLoop.py of Switch_Dict
instantiating app ryu.controller.ofp_handler of OFPHandler
```

pingall 结果

```

completed in 100.257 seconds
A ~/CS 1m 44s SDN_Spring_2024 18:28:31
> sudo mn --custom topo_1969_2.py --topo generated --controller remote
*** Creating network
*** Adding controller
Connecting to remote controller at 127.0.0.1:6653
*** Adding hosts:
SRI UCLA UCSB UTAH
*** Adding switches:
s1 s2 s3 s4
*** Adding links:
(s1, SRI) (s1, s2) (s1, s3) (s1, s4) (s2, UTAH) (s3, UCSB) (s3, s4) (s4, UCLA)
*** Configuring hosts
SRI UCLA UCSB UTAH
*** Starting controller
c0
*** Starting 4 switches
s1 s2 s3 s4 ...
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
SRI -> *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节, 来自 10.0.0.2: icmp_seq=1 ttl=64 时间=15.4 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 15.436/15.436/15.436/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节, 来自 10.0.0.3: icmp_seq=1 ttl=64 时间=6.73 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 6.730/6.730/6.730/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节, 来自 10.0.0.4: icmp_seq=1 ttl=64 时间=5.72 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 5.720/5.720/5.720/0.000 ms

X
UCLA -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节, 来自 10.0.0.1: icmp_seq=1 ttl=64 时间=0.251 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.251/0.251/0.251/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节, 来自 10.0.0.3: icmp_seq=1 ttl=64 时间=6.27 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 6.274/6.274/6.274/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节, 来自 10.0.0.4: icmp_seq=1 ttl=64 时间=5.73 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 5.733/5.733/5.733/0.000 ms

X
UCSB -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节, 来自 10.0.0.1: icmp_seq=1 ttl=64 时间=0.095 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.095/0.095/0.095/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节, 来自 10.0.0.2: icmp_seq=1 ttl=64 时间=0.067 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.067/0.067/0.067/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节, 来自 10.0.0.4: icmp_seq=1 ttl=64 时间=5.58 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 5.580/5.580/5.580/0.000 ms

X
UTAH -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节, 来自 10.0.0.1: icmp_seq=1 ttl=64 时间=0.036 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.036/0.036/0.036/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节, 来自 10.0.0.2: icmp_seq=1 ttl=64 时间=0.877 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.877/0.877/0.877/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节, 来自 10.0.0.3: icmp_seq=1 ttl=64 时间=0.615 毫秒

--- 10.0.0.3 ping 统计 ---

```



```
已发送 1 个包, 已接收 1 个包, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.615/0.615/0.615/0.000 ms

X
*** Results: 100% dropped (0/12 received)
mininet>
```

查看各Switch流表

```
sudo ovs-ofctl dump-flows s1 # SRI
sudo ovs-ofctl dump-flows s2 # UCLA
sudo ovs-ofctl dump-flows s3 # UCSB
sudo ovs-ofctl dump-flows s4 # UTAH
```

```
mininet> sh sudo ovs-ofctl dump-flows s1
cookie=0x0, duration=731.823s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s1-eth4",dl_dst=42:06:16:93:ae:65 actions=output:"s1-eth1"
cookie=0x0, duration=731.822s, table=0, n_packets=2, n_bytes=148, priority=10,in_port="s1-eth1",dl_dst=de:9a:83:4b:8a:47 actions=output:"s1-eth4"
cookie=0x0, duration=731.812s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s1-eth3",dl_dst=42:06:16:93:ae:65 actions=output:"s1-eth1"
cookie=0x0, duration=731.811s, table=0, n_packets=2, n_bytes=148, priority=10,in_port="s1-eth1",dl_dst=d2:50:34:4a:89:14 actions=output:"s1-eth3"
cookie=0x0, duration=731.803s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s1-eth2",dl_dst=42:06:16:93:ae:65 actions=output:"s1-eth1"
cookie=0x0, duration=731.802s, table=0, n_packets=2, n_bytes=148, priority=10,in_port="s1-eth1",dl_dst=9e:54:d4:3a:98:84 actions=output:"s1-eth2"
cookie=0x0, duration=731.782s, table=0, n_packets=4, n_bytes=336, priority=10,in_port="s1-eth2",dl_dst=de:9a:83:4b:8a:47 actions=output:"s1-eth4"
cookie=0x0, duration=731.781s, table=0, n_packets=5, n_bytes=378, priority=10,in_port="s1-eth4",dl_dst=9e:54:d4:3a:98:84 actions=output:"s1-eth2"
cookie=0x0, duration=731.780s, table=0, n_packets=5, n_bytes=378, priority=10,in_port="s1-eth3",dl_dst=9e:54:d4:3a:98:84 actions=output:"s1-eth2"
cookie=0x0, duration=731.767s, table=0, n_packets=4, n_bytes=336, priority=10,in_port="s1-eth2",dl_dst=d2:50:34:4a:89:14 actions=output:"s1-eth3"
cookie=0x0, duration=731.829s, table=0, n_packets=120, n_bytes=12228, priority=0 actions=CONTROLLER:65535
mininet> sh sudo ovs-ofctl dump-flows s2
cookie=0x0, duration=733.356s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s2-eth1",dl_dst=42:06:16:93:ae:65 actions=output:"s2-eth2"
cookie=0x0, duration=733.354s, table=0, n_packets=14, n_bytes=1892, priority=10,in_port="s2-eth2",dl_dst=9e:54:d4:3a:98:84 actions=output:"s2-eth1"
cookie=0x0, duration=733.355s, table=0, n_packets=4, n_bytes=336, priority=10,in_port="s2-eth1",dl_dst=de:9a:83:4b:8a:47 actions=output:"s2-eth2"
cookie=0x0, duration=733.328s, table=0, n_packets=4, n_bytes=336, priority=10,in_port="s2-eth1",dl_dst=d2:50:34:4a:89:14 actions=output:"s2-eth2"
cookie=0x0, duration=735.388s, table=0, n_packets=52, n_bytes=4832, priority=0 actions=CONTROLLER:65535
mininet> sh sudo ovs-ofctl dump-flows s3
cookie=0x0, duration=735.729s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s3-eth1",dl_dst=42:06:16:93:ae:65 actions=output:"s3-eth2"
cookie=0x0, duration=735.726s, table=0, n_packets=7, n_bytes=518, priority=10,in_port="s3-eth2",dl_dst=d2:50:34:4a:89:14 actions=output:"s3-eth1"
cookie=0x0, duration=735.709s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s3-eth1",dl_dst=de:9a:83:4b:8a:47 actions=output:"s3-eth3"
cookie=0x0, duration=735.705s, table=0, n_packets=2, n_bytes=148, priority=10,in_port="s3-eth2",dl_dst=d2:50:34:4a:89:14 actions=output:"s3-eth1"
cookie=0x0, duration=737.744s, table=0, n_packets=91, n_bytes=8838, priority=0 actions=CONTROLLER:65535
mininet> sh sudo ovs-ofctl dump-flows s4
cookie=0x0, duration=737.627s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s4-eth1",dl_dst=42:06:16:93:ae:65 actions=output:"s4-eth2"
cookie=0x0, duration=737.621s, table=0, n_packets=7, n_bytes=518, priority=10,in_port="s4-eth2",dl_dst=de:9a:83:4b:8a:47 actions=output:"s4-eth1"
cookie=0x0, duration=737.592s, table=0, n_packets=3, n_bytes=238, priority=10,in_port="s4-eth3",dl_dst=de:9a:83:4b:8a:47 actions=output:"s4-eth1"
cookie=0x0, duration=737.591s, table=0, n_packets=2, n_bytes=148, priority=10,in_port="s4-eth1",dl_dst=d2:50:34:4a:89:14 actions=output:"s4-eth3"
cookie=0x0, duration=739.638s, table=0, n_packets=92, n_bytes=8972, priority=0 actions=CONTROLLER:65535
mininet>
```

3. 附加题

实验任务二只给出了一种参考方案，SDN 中还有多种方案可供选择，请尝试设计实现一种新的策略解决上述环路广播问题

3.1 设计原理

我们给任一 ARP Record 赋予一定的生命周期，在这个时间范围内：

- 如果出现了多个 ARP Request 报文，则只记录第一个报文，忽略后面的报文

查询资料可知：一个 ARP Record 的生存期大约为 1min (60s)

3.2 代码实现

task3_arp_sdn.py:

PYTHON

```
from ryu.base import app_manager
from ryu.controller import ofp_event
from ryu.controller.handler import MAIN_DISPATCHER, CONFIG_DISPATCHER
from ryu.controller.handler import set_ev_cls
from ryu.ofproto import ofproto_v1_3
from ryu.lib.packet import packet
from ryu.lib.packet import ethernet
from ryu.lib.packet import arp
from ryu.lib.packet import ether_types

ETHERNET = ethernet.ethernet.__name__
ETHERNET_MULTICAST = "ff:ff:ff:ff:ff:ff"
ARP = arp.arp.__name__
ARP_TIMEOUT = 60

class Switch_Dict(app_manager.RyuApp):
    OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]

    def __init__(self, *args, **kwargs):
        super(Switch_Dict, self).__init__(*args, **kwargs)
        self.mac_to_port = {}

        # recording the latest TimeNode (a ARP Request Packet coming in)
        self.latest_stamp = {} # dpid, src_ip, dst_ip -> timestamp

    def add_flow(self, datapath, priority, match, actions, idle_timeout=0, h
        dp = datapath
        ofp = dp.ofproto
        parser = dp.ofproto_parser

        inst = [parser.OFPInstructionActions(ofp.OFPIT_APPLY_ACTIONS, actions
        mod = parser.OFPFlowMod(
            datapath=dp,
            priority=priority,
            idle_timeout=idle_timeout,
            hard_timeout=hard_timeout,
            match=match,
            instructions=inst,
        )
        dp.send_msg(mod)
```

```

@set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
def switch_features_handler(self, ev):
    msg = ev.msg
    dp = msg.datapath
    ofp = dp.ofproto
    parser = dp.ofproto_parser

    '''
    创建一个动作列表，表示将数据包发送到控制器：
    - ofp.OFPP_CONTROLLER表示数据包将发送到控制器， ofp.OFPCML_NO_BUFFER表示不
    '''

    match = parser.OFPMatch()
    actions = [parser.OFPACTIONOutput(ofp.OFPP_CONTROLLER, ofp.OFPCML_NO_
    '''

    调用add_flow方法向交换机添加流表规则：
    - 这条规则的优先级为0，表示最低优先级，匹配所有数据包，并将它们发送到控制器
    '''

    self.add_flow(dp, 0, match, actions)

    # let switch send arp to controller
    '''
    创建一个匹配条件，匹配目的MAC地址为广播地址（ETHERNET_MULTICAST）的数据包。
    '''

    match = parser.OFPMatch(eth_dst=ETHERNET_MULTICAST)
    actions = [parser.OFPACTIONOutput(ofp.OFPP_CONTROLLER, ofp.OFPCML_NO_

    # a little higher priority to make switch must send arp to controller
    '''

    确保交换机会将ARP请求数据包发送到控制器
    '''

    self.add_flow(dp, 1, match, actions)

# what we actually done
@set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
def packet_in_handler(self, ev):
    msg = ev.msg
    dp = msg.datapath
    ofp = dp.ofproto
    parser = dp.ofproto_parser

    # the identity of switch
    dpid = dp.id
    self.mac_to_port.setdefault(dpid, {})
    self.latest_stamp.setdefault(dpid, {})

```

```

# the port that receive the packet
in_port = msg.match["in_port"]
pkt = packet.Packet(msg.data)
eth_pkt = pkt.get_protocol(ethernet.ethernet)
if eth_pkt.ethertype == ether_types.ETH_TYPE_LLDP:
    return
if eth_pkt.ethertype == ether_types.ETH_TYPE_IPV6:
    return

# get the mac
eth_dst = eth_pkt.dst
eth_src = eth_pkt.src

self.logger.info(
    "Pkt-in: SW(%s) Src(%s) Dst(%s) InPort(%s)", dpid, eth_src, eth_d
)

if eth_src not in self.mac_to_port[dpid]:
    # recording srcMAC and mapping(arcMAC, portIn)
    self.mac_to_port[dpid].setdefault(eth_src, {})
    self.mac_to_port[dpid][eth_src] = in_port

# ARP request
arp_pkt = pkt.get_protocol(arp.arp)
currentTime = ev.timestamp

if arp_pkt and arp_pkt.opcode == arp.ARP_REQUEST:
    # get the ip
    arp_src_ip = arp_pkt.src_ip
    arp_dst_ip = arp_pkt.dst_ip

    # arp for the first time
    if arp_src_ip not in self.latest_stamp[dpid]:
        self.latest_stamp[dpid].setdefault(arp_src_ip, {})
        self.latest_stamp[dpid][arp_src_ip][arp_dst_ip] = currentTime
    # another dst ip, update, too
    elif arp_dst_ip not in self.latest_stamp[dpid][arp_src_ip]:
        self.latest_stamp[dpid][arp_src_ip].setdefault(arp_dst_ip, {})
        self.latest_stamp[dpid][arp_src_ip][arp_dst_ip] = currentTime
    # arp for the second time
    elif currentTime - self.latest_stamp[dpid][arp_src_ip][arp_dst_ip] > 0:
        print("duration is not long enough!")
        return

# self-learning
if eth_dst in self.mac_to_port[dpid]:

```

```
        out_port = self.mac_to_port[dpid][eth_dst]
    else:
        out_port = ofp.OFPP_FLOOD

    actions = [parser.OFPActionOutput(out_port)]

    if out_port != ofp.OFPP_FLOOD:
        match = parser.OFPMatch(in_port=in_port, eth_dst=eth_dst)
        self.add_flow(dp, 10, match, actions)

    data = None
    if msg.buffer_id == ofp.OFP_NO_BUFFER:
        data = msg.data
    out = parser.OFPPacketOut(
        datapath=dp, buffer_id=msg.buffer_id, in_port=in_port, actions=ac
    )
    dp.send_msg(out)
```

3.3 实验分析

安装控制器并进行对接

```
ryu-manager --observe-links task3_arp_sdn.py
```



```

A@~ /CS
$ sudo mn --custom topo_1962_2.py --topo generated --controller remote
[sudo] root-hbx 的密码:
*** Creating network
*** Adding controller
Connecting to remote controller at 127.0.0.1:6653
*** Adding hosts:
SRI UCLA UCSB UTAH
*** Adding switches:
s1 s2 s3 s4
*** Adding links:
(s1, SRI) (s1, s2) (s1, s3) (s1, s4) (s2, UTAH) (s3, UCSB) (s3, s4) (s4, UCLA)
*** Configuring hosts
SRI UCLA UCSB UTAH
*** Starting controller
c0
*** Starting 4 switches
s1 s2 s3 s4 ...
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
SRI -> *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节，来自 10.0.0.2: icmp_seq=1 ttl=64 时间=11.3 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 11.289/11.289/11.289/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节，来自 10.0.0.3: icmp_seq=1 ttl=64 时间=6.49 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 6.490/6.490/6.490/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节，来自 10.0.0.4: icmp_seq=1 ttl=64 时间=5.45 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 5.452/5.452/5.452/0.000 ms

X
UCLA -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节，来自 10.0.0.1: icmp_seq=1 ttl=64 时间=0.047 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 0.047/0.047/0.047/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节，来自 10.0.0.3: icmp_seq=1 ttl=64 时间=9.95 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 9.945/9.945/9.945/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节，来自 10.0.0.4: icmp_seq=1 ttl=64 时间=10.3 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 10.258/10.258/10.258/0.000 ms

X
UCSB -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节，来自 10.0.0.1: icmp_seq=1 ttl=64 时间=0.051 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 0.051/0.051/0.051/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节，来自 10.0.0.2: icmp_seq=1 ttl=64 时间=0.034 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 0.034/0.034/0.034/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.4 (10.0.0.4) 56(84) 字节的数据。
64 字节，来自 10.0.0.4: icmp_seq=1 ttl=64 时间=6.65 毫秒

--- 10.0.0.4 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 6.645/6.645/6.645/0.000 ms

X
UTAH -> *** Error: could not parse ping output: PING 10.0.0.1 (10.0.0.1) 56(84) 字节的数据。
64 字节，来自 10.0.0.1: icmp_seq=1 ttl=64 时间=0.031 毫秒

--- 10.0.0.1 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 0.031/0.031/0.031/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.2 (10.0.0.2) 56(84) 字节的数据。
64 字节，来自 10.0.0.2: icmp_seq=1 ttl=64 时间=1.03 毫秒

--- 10.0.0.2 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 1.030/1.030/1.030/0.000 ms

X *** Error: could not parse ping output: PING 10.0.0.3 (10.0.0.3) 56(84) 字节的数据。
64 字节，来自 10.0.0.3: icmp_seq=1 ttl=64 时间=1.00 毫秒

--- 10.0.0.3 ping 统计 ---
已发送 1 个包，已接收 1 个包，0% packet loss, time 0ms
rtt min/avg/max/ndev = 1.004/1.004/1.004/0.000 ms

X
*** Results: 100% dropped (0/12 received)
mininet>

```

查看各Switch流表

```

sudo ovs-ofctl dump-flows s1 # SRI
sudo ovs-ofctl dump-flows s2 # UCLA
sudo ovs-ofctl dump-flows s3 # UCSB
sudo ovs-ofctl dump-flows s4 # UTAH

```

```

mininet> sh sudo ovs-ofctl dump-flows s1
cookie=0x0, duration=177.826s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s1-eth4", dl_dst=ca:84:a3:a1:56:6e actions=output:"s1-eth1"
cookie=0x0, duration=177.825s, table=0, n_packets=2, n_bytes=148, priority=18, in_port="s1-eth1", dl_dst=ad:33:a2:87:c3 actions=output:"s1-eth4"
cookie=0x0, duration=177.814s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s1-eth3", dl_dst=ca:84:a3:a1:56:6e actions=output:"s1-eth1"
cookie=0x0, duration=177.814s, table=0, n_packets=2, n_bytes=148, priority=18, in_port="s1-eth1", dl_dst=a6:c3:38:84:8f:a4 actions=output:"s1-eth3"
cookie=0x0, duration=177.806s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s1-eth2", dl_dst=ca:84:a3:a1:56:6e actions=output:"s1-eth1"
cookie=0x0, duration=177.805s, table=0, n_packets=2, n_bytes=148, priority=18, in_port="s1-eth1", dl_dst=7a:b6:e5:1c:63:1a actions=output:"s1-eth2"
cookie=0x0, duration=177.777s, table=0, n_packets=4, n_bytes=236, priority=18, in_port="s1-eth2", dl_dst=16:ad:33:a2:87:c3 actions=output:"s1-eth4"
cookie=0x0, duration=177.773s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s1-eth4", dl_dst=7a:b6:e5:1c:63:1a actions=output:"s1-eth2"
cookie=0x0, duration=177.778s, table=0, n_packets=4, n_bytes=236, priority=18, in_port="s1-eth3", dl_dst=7a:b6:e5:1c:63:1a actions=output:"s1-eth2"
cookie=0x0, duration=177.758s, table=0, n_packets=4, n_bytes=236, priority=18, in_port="s1-eth2", dl_dst=a6:c3:38:84:8f:a4 actions=output:"s1-eth3"
cookie=0x0, duration=181.198s, table=0, n_packets=3, n_bytes=238, priority=1, dl_dst=ff:ff:ff:ff:ff:ff actions=CONTROLLER:65535
cookie=0x0, duration=181.198s, table=0, n_packets=97, n_bytes=18648, priority=0 actions=CONTROLLER:65535
mininet> sh sudo ovs-ofctl dump-flows s2
cookie=0x0, duration=188.898s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s2-eth1", dl_dst=ca:84:a3:a1:56:6e actions=output:"s2-eth2"
cookie=0x0, duration=188.898s, table=0, n_packets=14, n_bytes=1892, priority=18, in_port="s2-eth2", dl_dst=7a:b6:e5:1c:63:1a actions=output:"s2-eth1"
cookie=0x0, duration=188.871s, table=0, n_packets=4, n_bytes=336, priority=18, in_port="s2-eth1", dl_dst=16:ad:33:a2:87:c3 actions=output:"s2-eth2"
cookie=0x0, duration=188.852s, table=0, n_packets=4, n_bytes=336, priority=18, in_port="s2-eth1", dl_dst=a6:c3:38:84:8f:a4 actions=output:"s2-eth2"
cookie=0x0, duration=184.283s, table=0, n_packets=6, n_bytes=252, priority=1, dl_dst=ff:ff:ff:ff:ff:ff actions=CONTROLLER:65535
cookie=0x0, duration=184.283s, table=0, n_packets=48, n_bytes=4808, priority=0 actions=CONTROLLER:65535
mininet> sh sudo ovs-ofctl dump-flows s3
cookie=0x0, duration=186.569s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s3-eth1", dl_dst=ca:84:a3:a1:56:6e actions=output:"s3-eth2"
cookie=0x0, duration=186.567s, table=0, n_packets=3, n_bytes=518, priority=18, in_port="s3-eth2", dl_dst=a6:c3:38:84:8f:a4 actions=output:"s3-eth1"
cookie=0x0, duration=186.545s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s3-eth1", dl_dst=16:ad:33:a2:87:c3 actions=output:"s3-eth3"
cookie=0x0, duration=186.542s, table=0, n_packets=2, n_bytes=148, priority=18, in_port="s3-eth3", dl_dst=a6:c3:38:84:8f:a4 actions=output:"s3-eth1"
cookie=0x0, duration=189.944s, table=0, n_packets=11, n_bytes=462, priority=1, dl_dst=ff:ff:ff:ff:ff:ff actions=CONTROLLER:65535
cookie=0x0, duration=189.944s, table=0, n_packets=72, n_bytes=7784, priority=0 actions=CONTROLLER:65535
mininet> sh sudo ovs-ofctl dump-flows s4
// cookie=0x0, duration=189.162s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s4-eth1", dl_dst=ca:84:a3:a1:56:6e actions=output:"s4-eth2"
cookie=0x0, duration=189.158s, table=0, n_packets=7, n_bytes=518, priority=18, in_port="s4-eth2", dl_dst=16:ad:33:a2:87:c3 actions=output:"s4-eth1"
cookie=0x0, duration=189.124s, table=0, n_packets=3, n_bytes=238, priority=18, in_port="s4-eth3", dl_dst=16:ad:33:a2:87:c3 actions=output:"s4-eth1"
cookie=0x0, duration=189.123s, table=0, n_packets=2, n_bytes=148, priority=18, in_port="s4-eth1", dl_dst=a6:c3:38:84:8f:a4 actions=output:"s4-eth3"
cookie=0x0, duration=192.524s, table=0, n_packets=18, n_bytes=428, priority=1, dl_dst=ff:ff:ff:ff:ff:ff actions=CONTROLLER:65535
cookie=0x0, duration=192.524s, table=0, n_packets=72, n_bytes=7784, priority=0 actions=CONTROLLER:65535
mininet>

```